

Solving problems by Searching for Solutions

AIMA Chapter 3.1 – 3.4

So far ...

What is AI?

What are agents?

- Rational agents

Task environment

- Specifying
- Properties

Types of agents

- Reflex (simple & model based)
- Goal based
- Utility based

Searching for Solutions

Agent knows the
current state

There are finite
number of actions

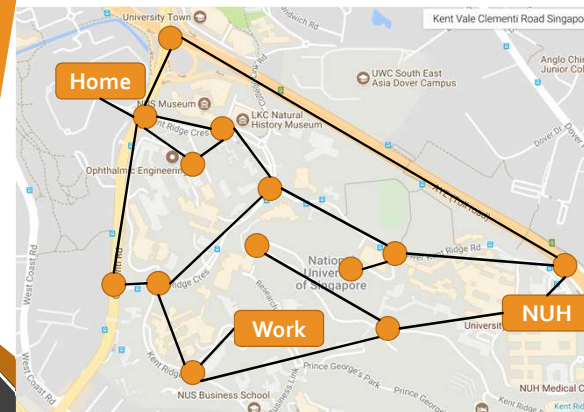
- Fully observable, deterministic, discrete environment

Action have exactly
one outcome

Examples

- Shortest route going through all cities
- Winning sequence of moves in a **single player** game
- Assembling a machine.
- Find path through a maze
- Route planning
- Solution to Boolean satisfiability problem

Example: Route Planning



Problem Formulation

Initial state

- The start state

Actions

- $ACTION(s)$: set of actions applicable in s ;
e.g., $GO(NUH)$

Transition model

- Description of what each action does
- $s' = RESULT(s, a)$ state that results from doing a in s

Problem Formulation

Goal Test

- Is the state s equal the goal state?
- Explicit set of goal states, e.g., $\{In(Work)\}$
- Implicit function, e.g., $IsCheckmate(s)$

Path Cost

- Additive. e.g., sum of distances, number of actions executed
- $c(s, a, s')$: the **step cost** of taking action a in state s to reach state s' .
Assumed to be ≥ 0


Solution: sequence of actions
leading from initial to goal state

Is the solution unique?

Is it optimal?

Can it be efficiently found?

Example: Route Planning



Example: 8 Puzzle

7	2	4
5		6
8	3	1


Start State

	1	2
3	4	5
6	7	8

Goal State

1. States (Initial State)?
2. Actions?
3. Transition Model?
4. Goal Test?
5. Cost Function?

Example: Item Assembly



1. States (Initial State)?
2. Actions?
3. Transition Model?
4. Goal Test?
5. Cost Function?

Tree Search Algorithms

- Start at the **source node**, keep searching until we reach a **goal state**.
- **Frontier**: nodes that we have seen but haven't explored yet (at initialization: the frontier is just the source)
- At each iteration, choose a node from the frontier, explore it, and **add its neighbors** to the frontier.

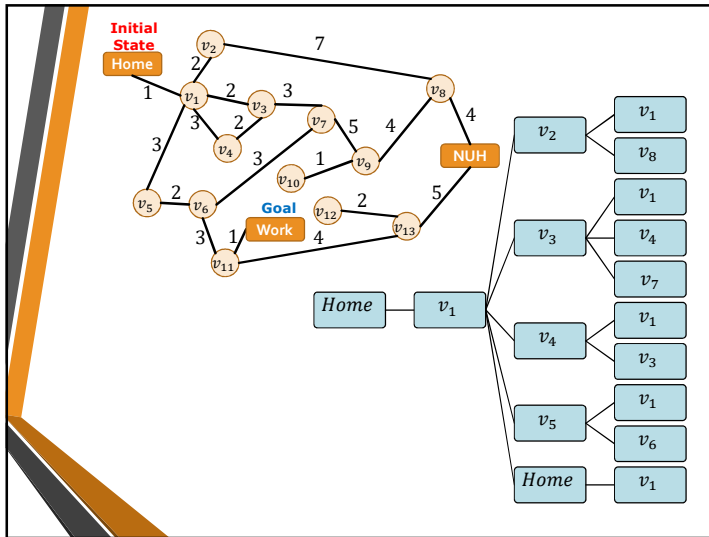
Tree Search Algorithms

- Start at the **source node**, keep searching until we

```

function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier
  
```

- At each iteration, choose a node from the frontier, explore it, and **add its neighbors** to the frontier.



Traversal Methods

- In what order do we explore the frontier nodes?
Search strategy
- Which neighbors do we add? All of them?

"Algorithms which do not remember their (search) history are doomed to repeat it"

Graph Search

- Graph Search "Remembers"!
- A node that's been explored once will not be revisited.

Graph Search

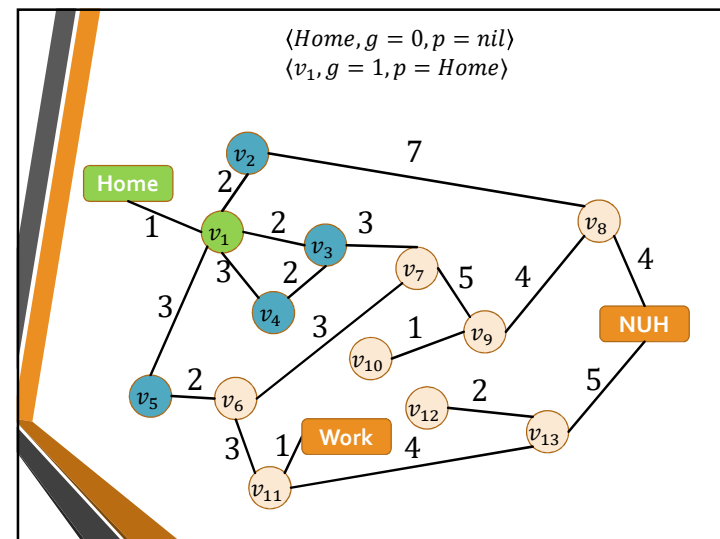
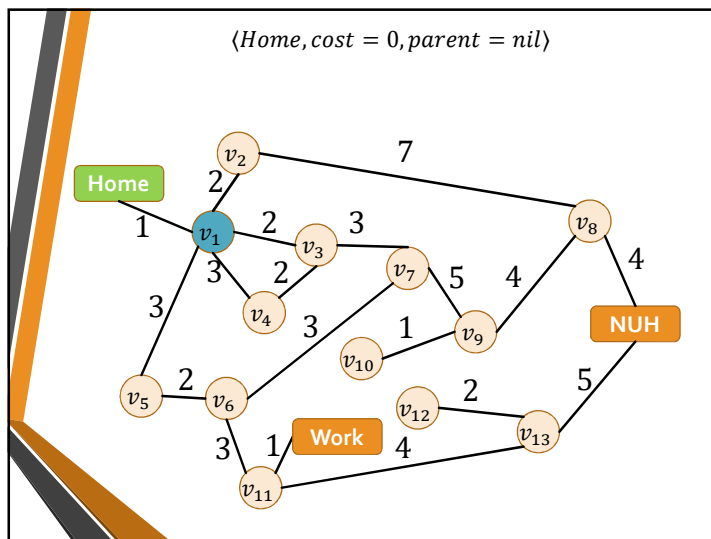
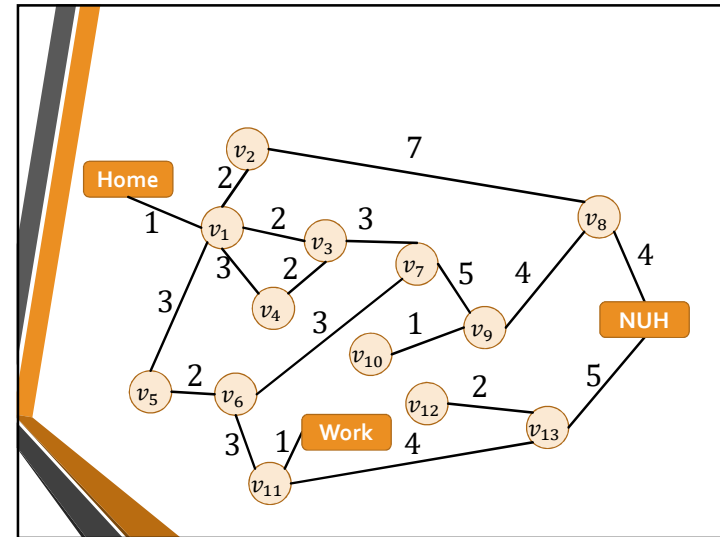

```

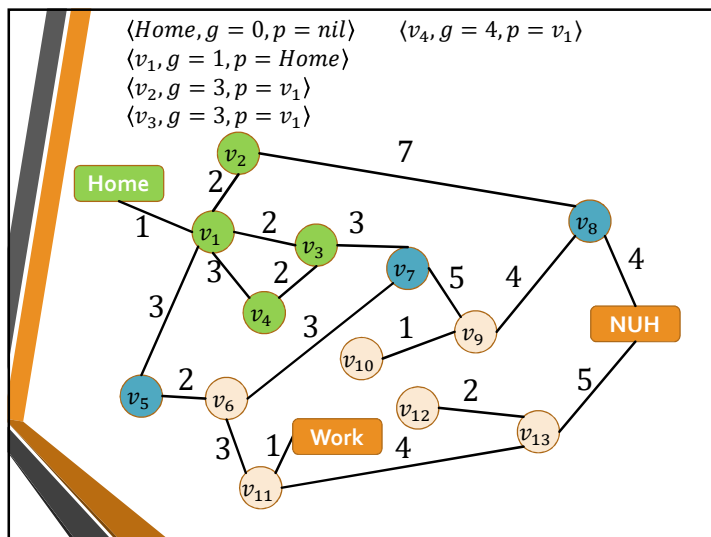
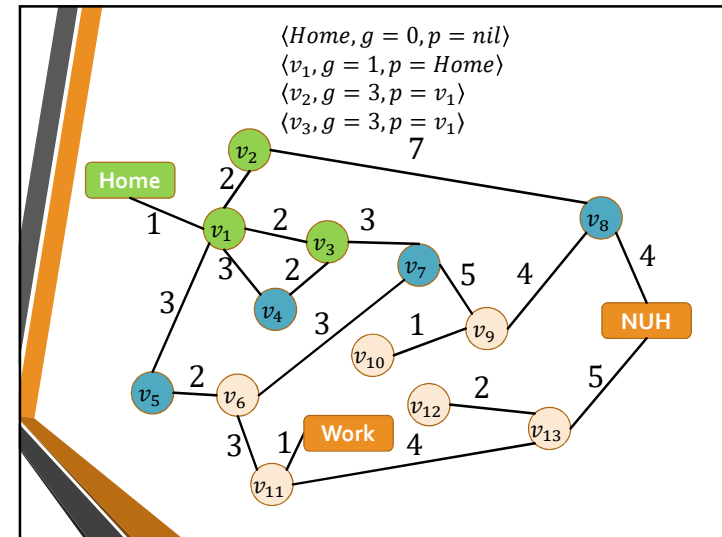
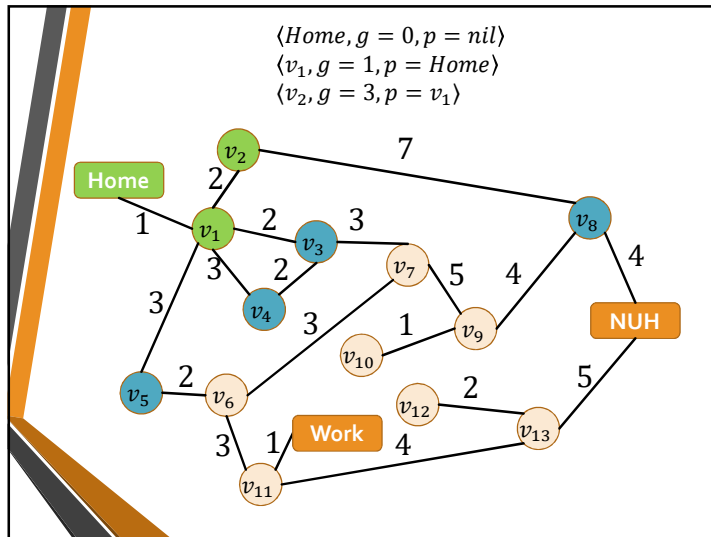
function GRAPH-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
    only if not in the frontier or explored set
  
```

The names 'Graph Search' and 'Tree Search' can be unintuitive: remember their definitions!

Implementation: States vs. Nodes

- A **state** represents a physical configuration
- A **node** is a data structure constituting part of search tree. It includes **state**, **parent node**, **action**, and **path cost** $g(n)$.
- Two different nodes are allowed to contain same world state





Search Strategies

Any deterministic search algorithm is determined by the order of node expansion.

Evaluation Criteria

- completeness: always find a solution if exists
- optimality: find a least-cost solution
- time complexity: number of nodes generated
- space complexity: max. number of nodes in memory

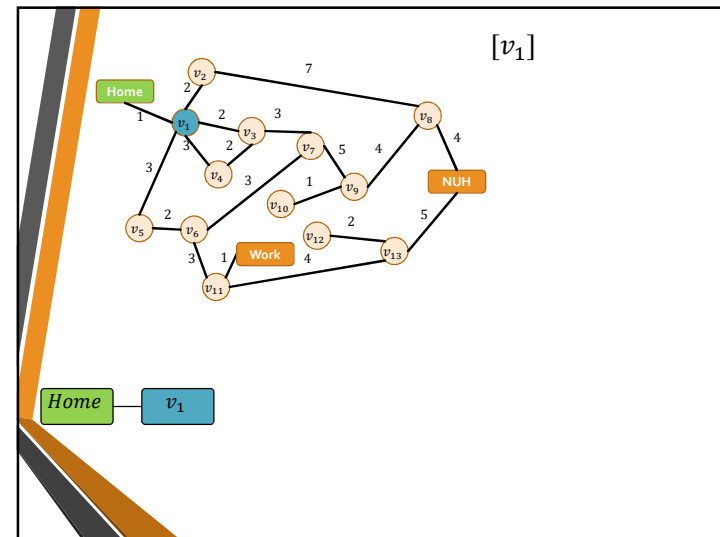
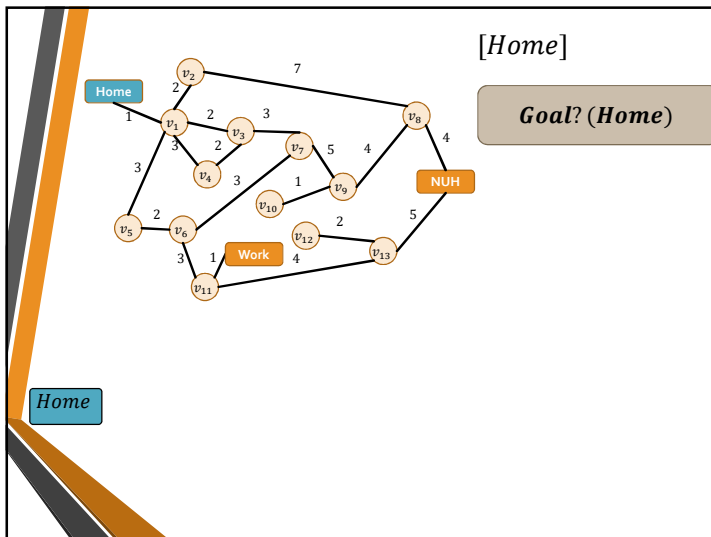
Problem Parameters

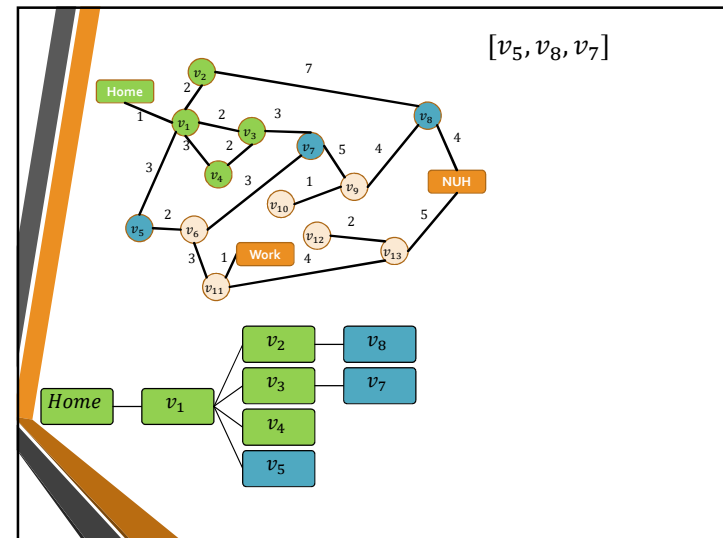
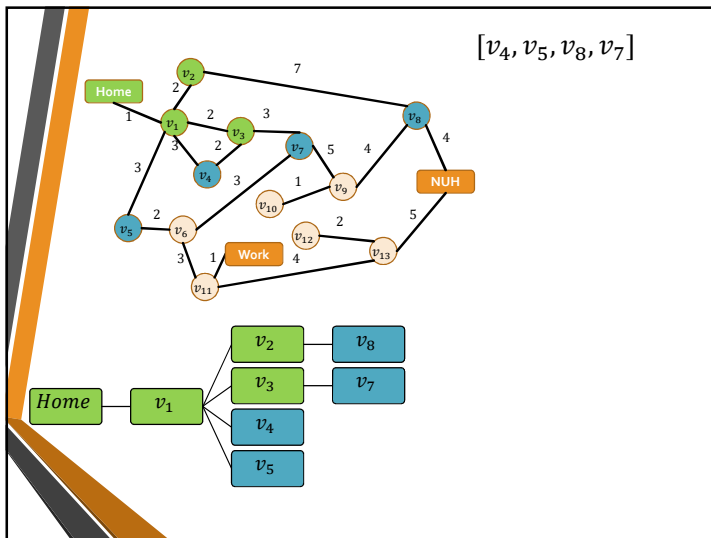
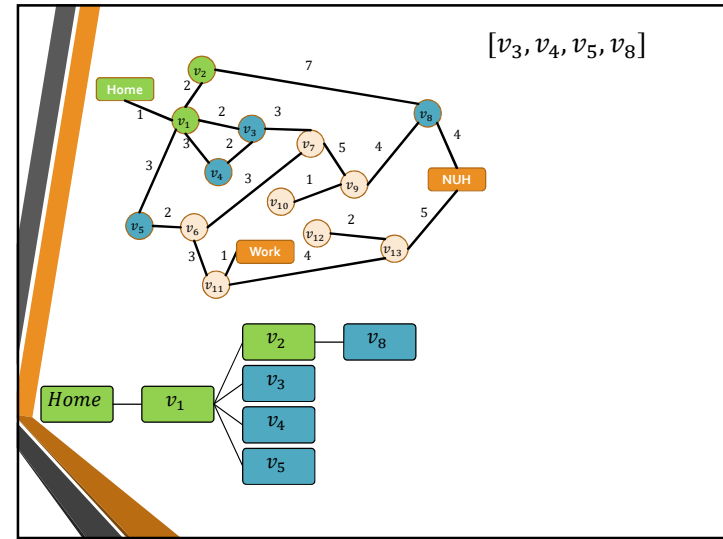
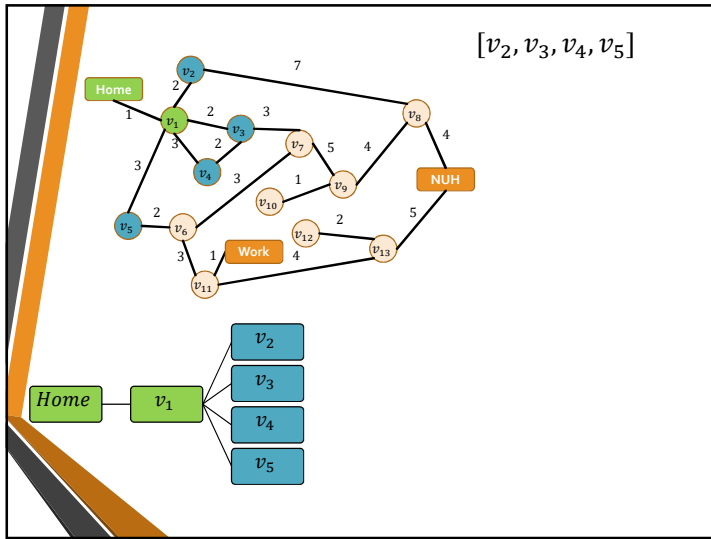
- b : maximum # of successors of any node (may be ∞)
- d : depth of shallowest **goal** node
- m : maximum depth of search tree (may be ∞)

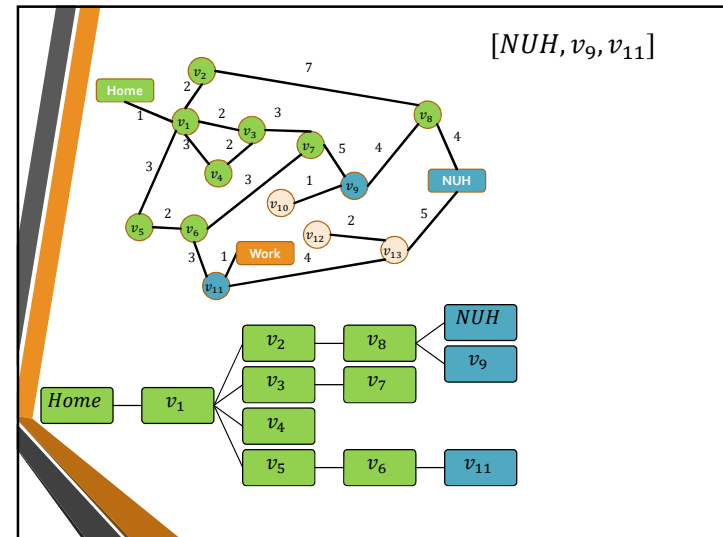
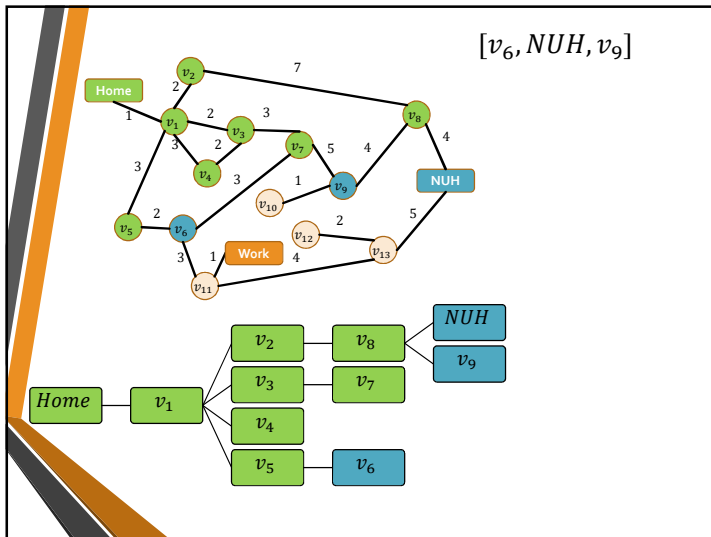
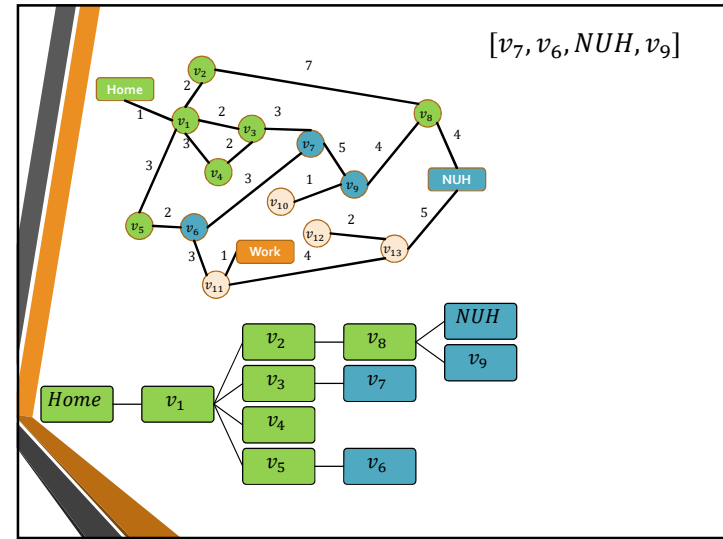
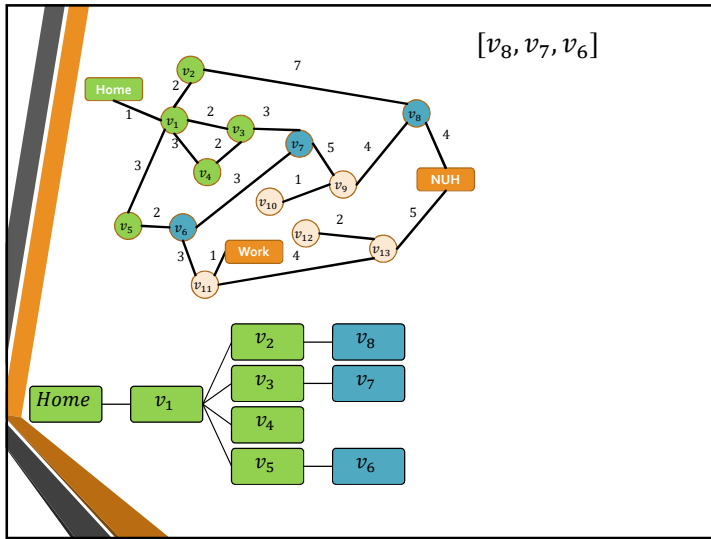


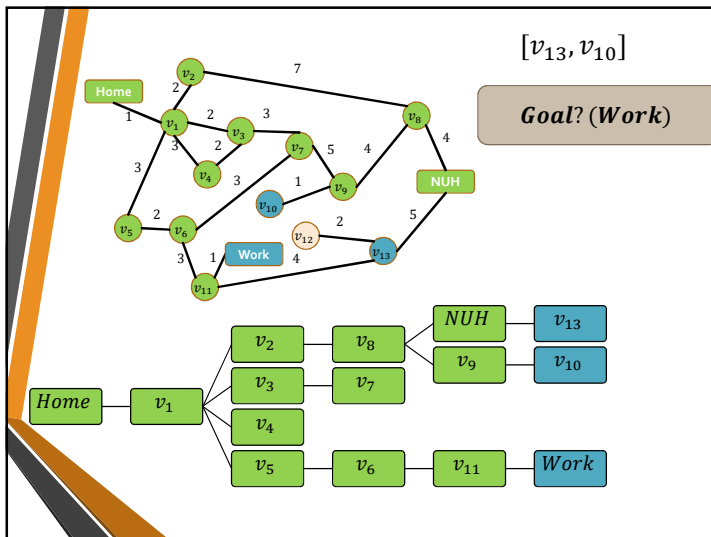
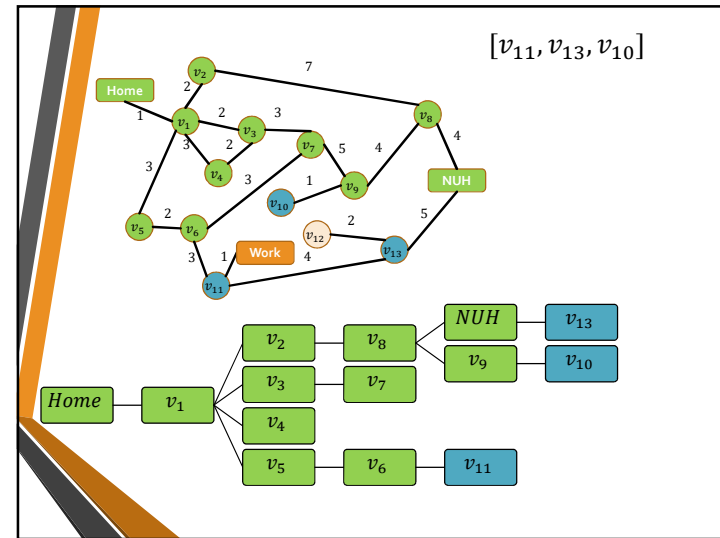
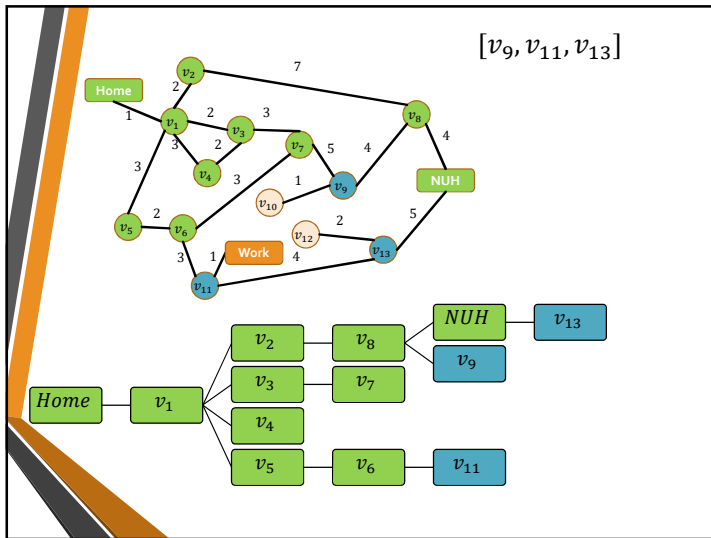
Breadth-First Search (BFS)

- **Idea:** Expand shallowest unexpanded node
- **Implementation:** Frontier is a FIFO queue, i.e., insert new successors at the end










Properties of BFS

Property	
Complete?	Yes (if b is finite)
Optimal	No (unless step cost = 1)
Time	$\mathcal{O}(b) + \mathcal{O}(b^2) + \dots + \mathcal{O}(b^d) = \mathcal{O}(b^d)$
Space	Max size of frontier $\mathcal{O}(b^d)$

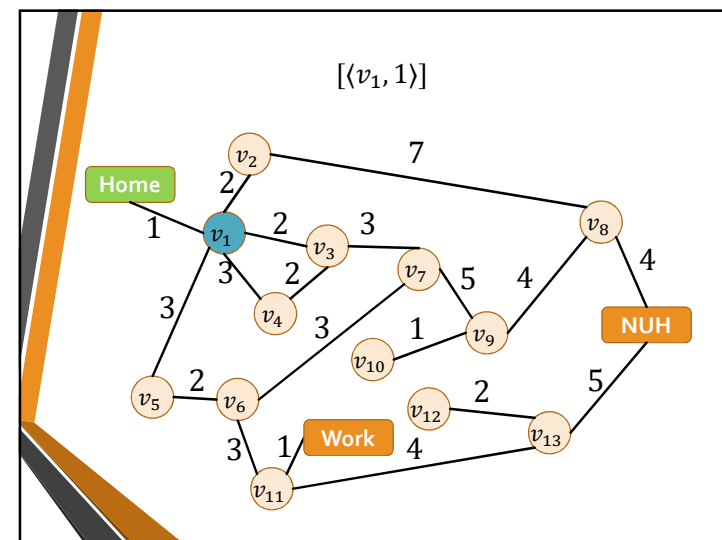
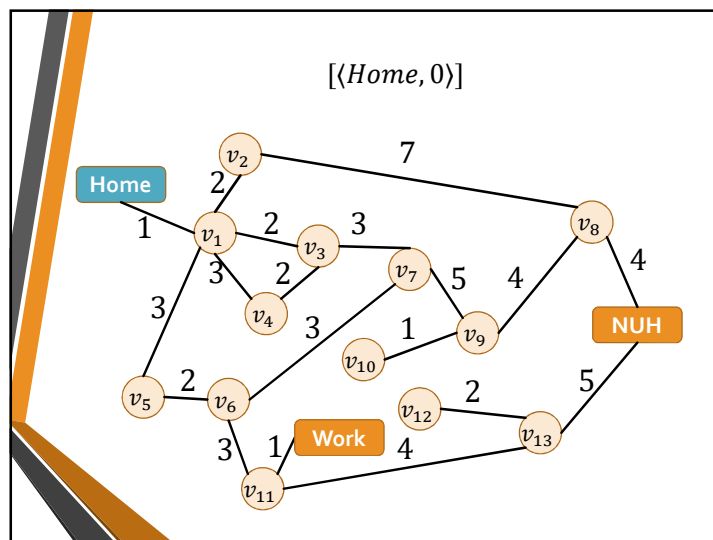
Space is the bigger problem (more than time)

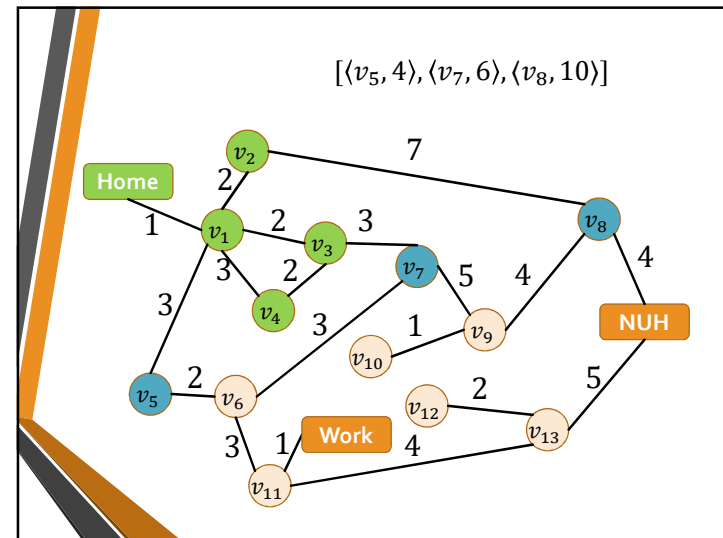
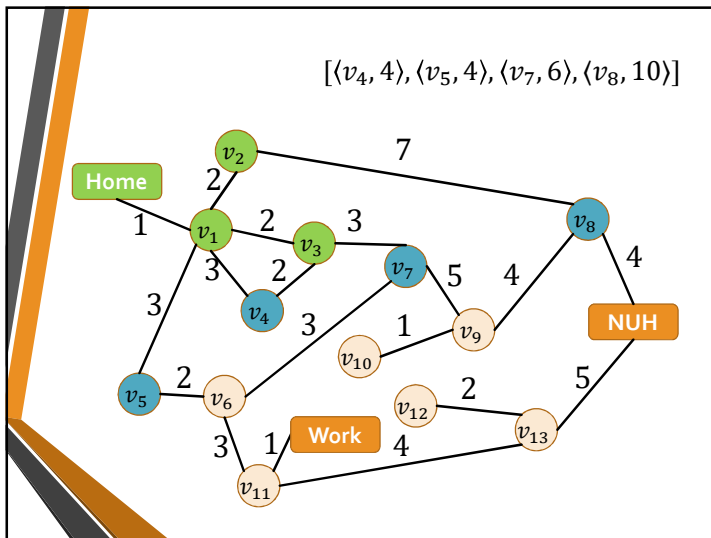
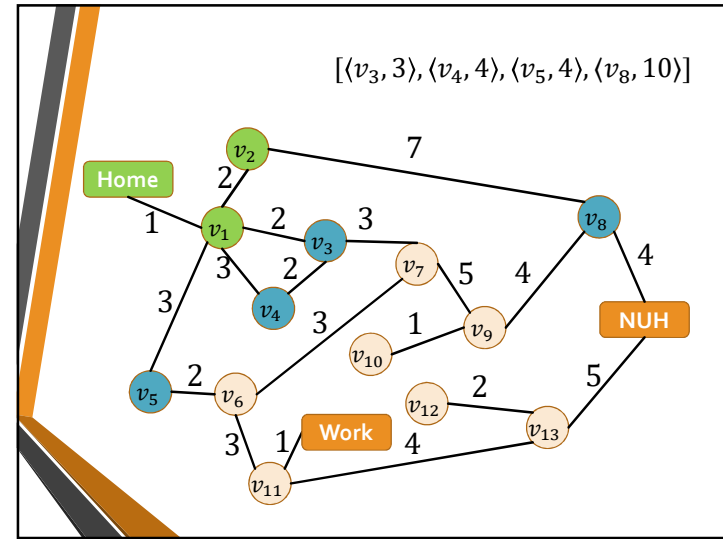
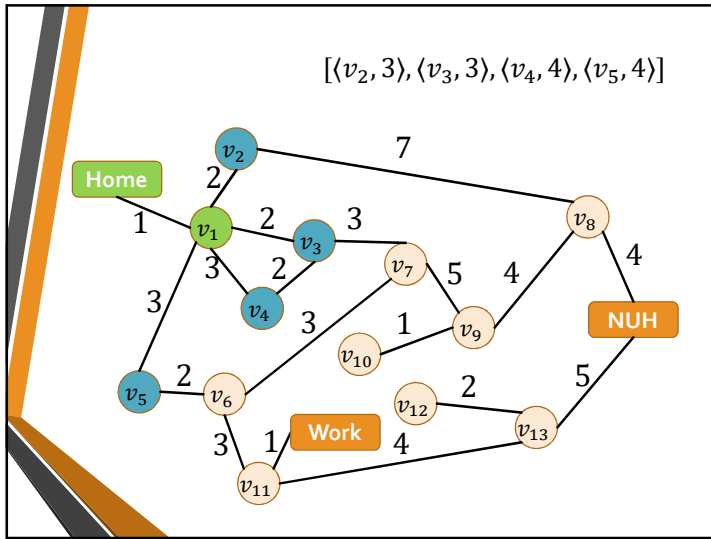
completeness: always find a solution if exists
 optimality: find a least-cost solution
 time complexity: number of nodes generated
 space complexity: max. number of nodes in memory

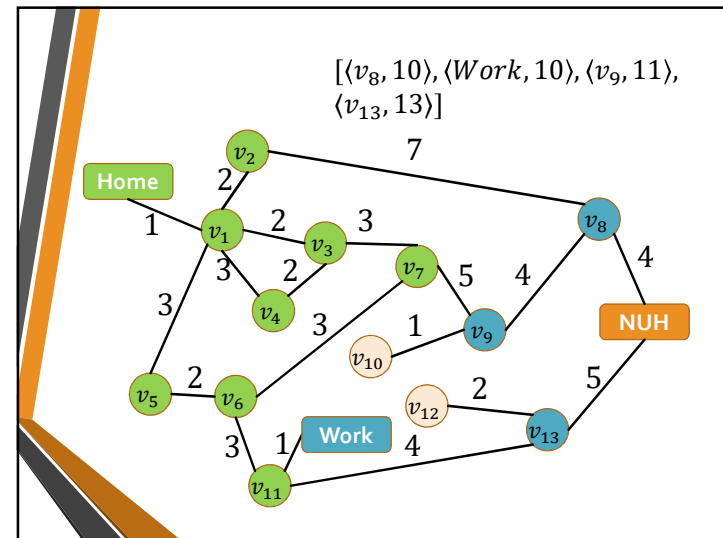
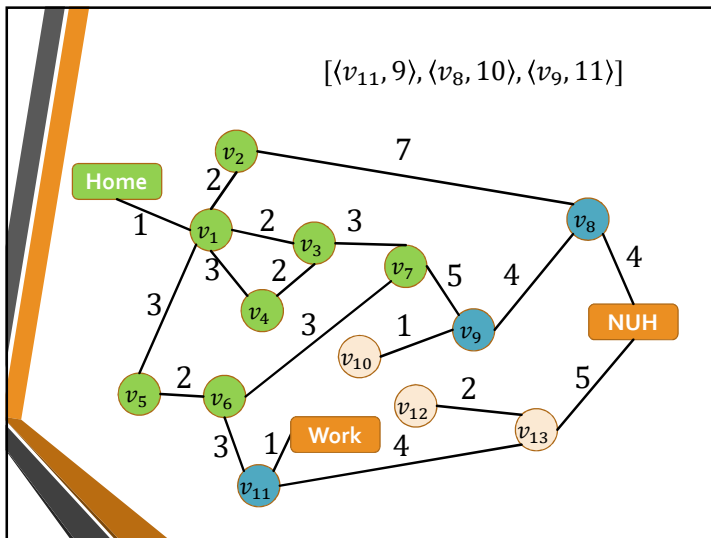
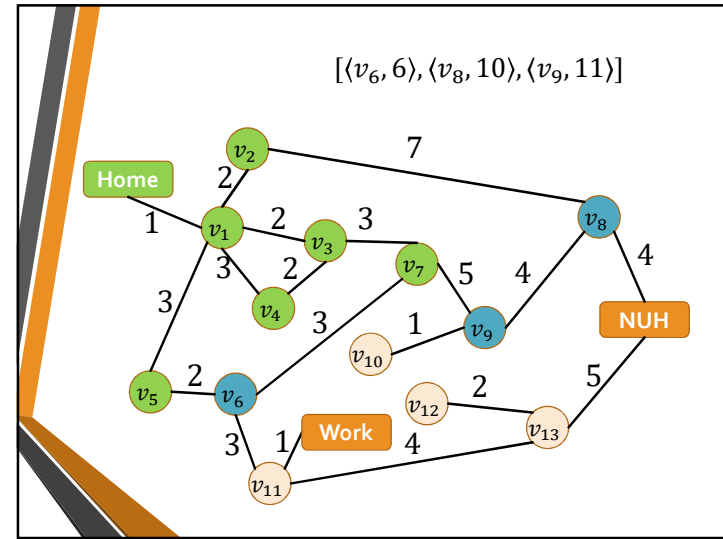
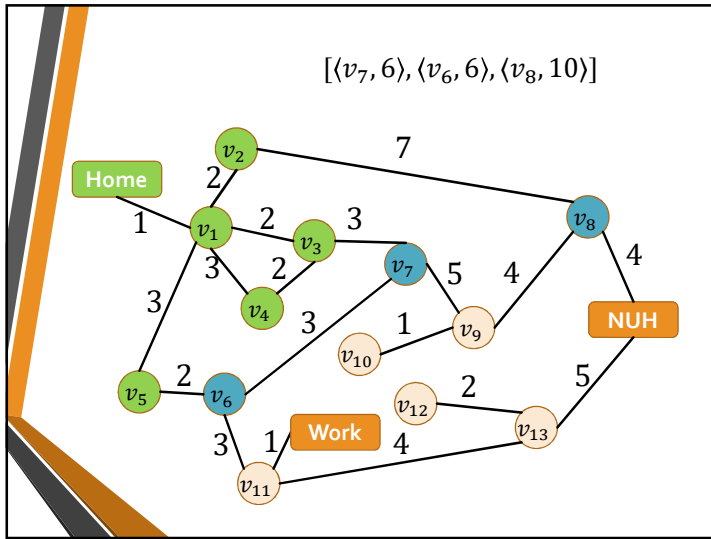


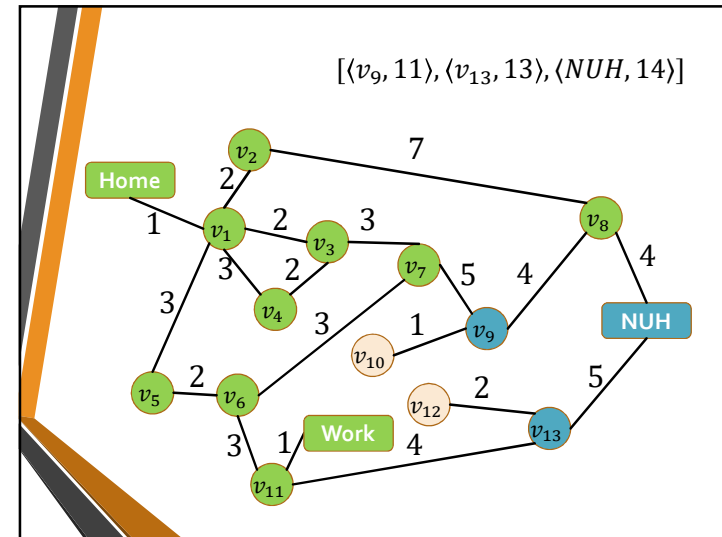
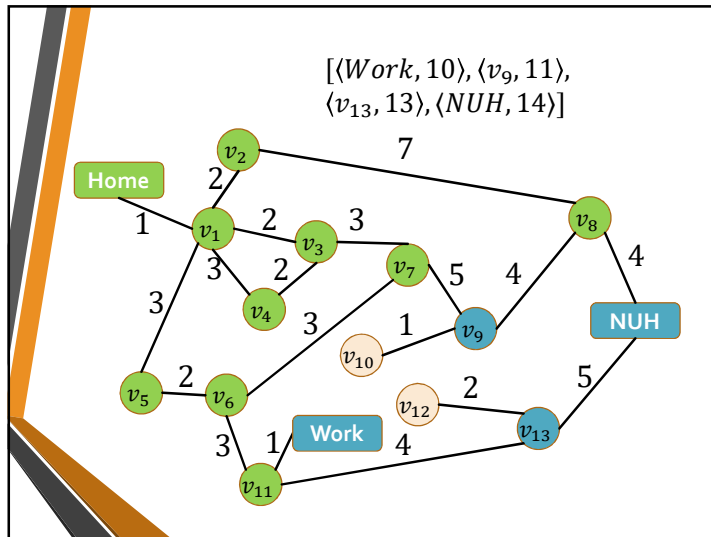
Uniform-Cost Search (UCS)

- Idea: expand least-path-cost unexpanded node
- Frontier: priority queue ordered by path cost g
- Equivalent to BFS if all step costs are equal









Properties of Uniform Cost Search

Property	
Complete?	Yes (if all step costs are $\geq \epsilon$)
Optimal	Yes (shortest path nodes expanded first)
Time	$\mathcal{O}(b^{1+\lceil \frac{C^*}{\epsilon} \rceil})$ where C^* is the optimal cost.
Space	$\mathcal{O}(b^{1+\lceil \frac{C^*}{\epsilon} \rceil})$

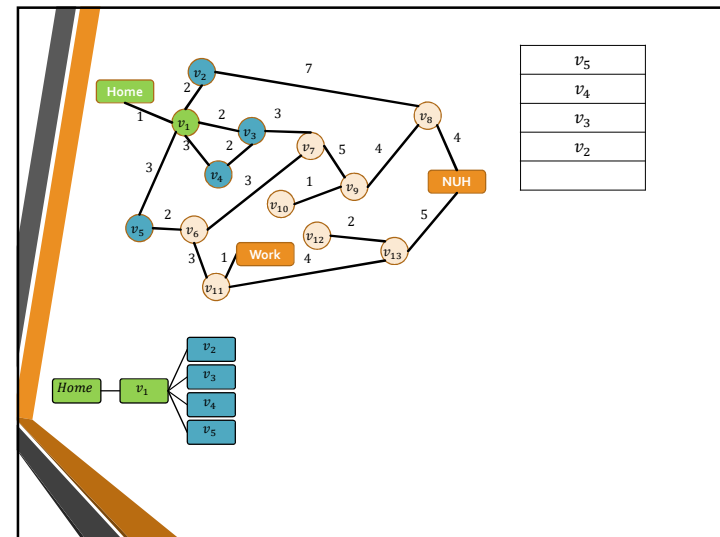
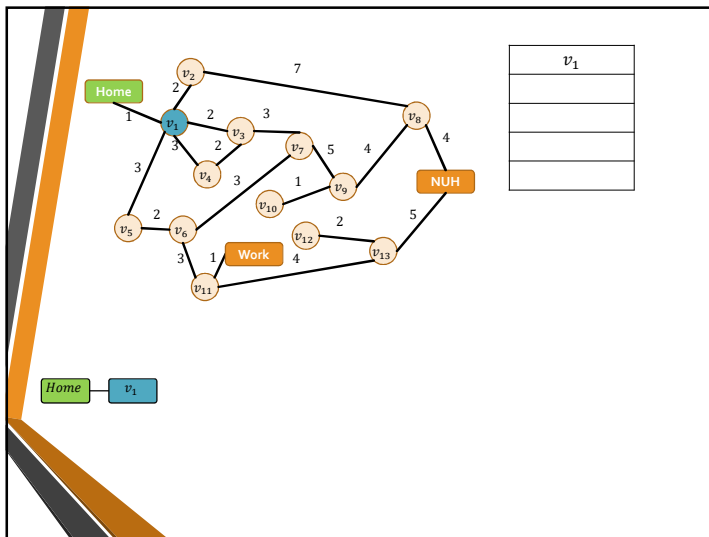
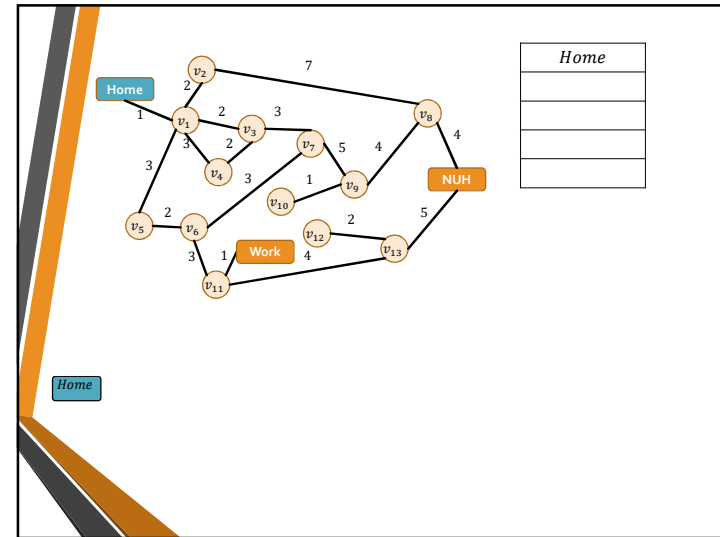
completeness: always find a solution if exists
 optimality: find a least-cost solution
 time complexity: number of nodes generated
 space complexity: max. number of nodes in memory

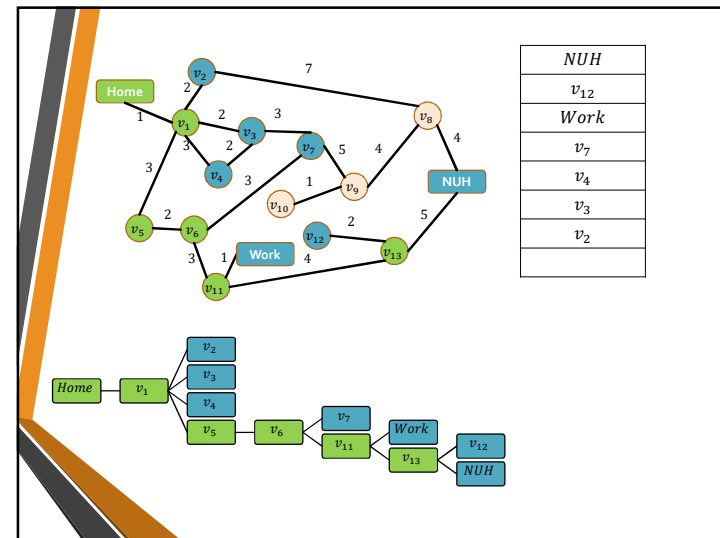
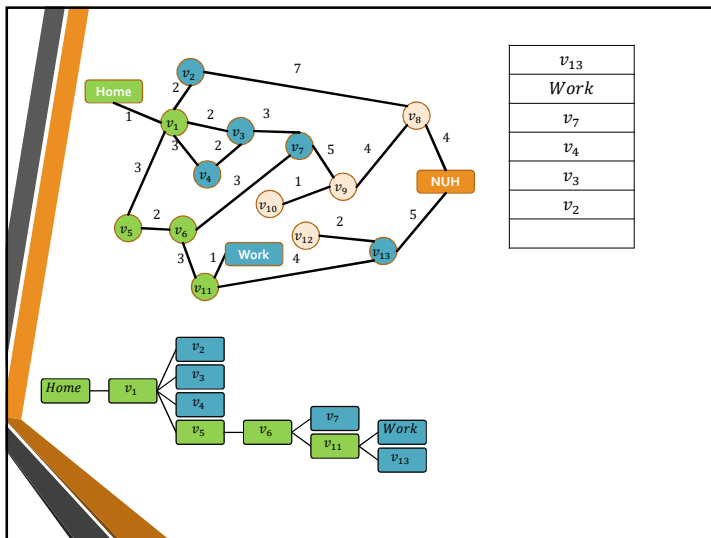
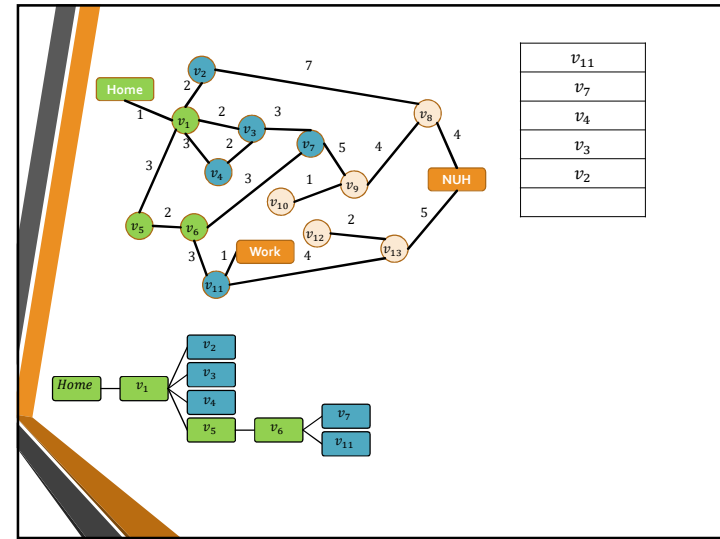
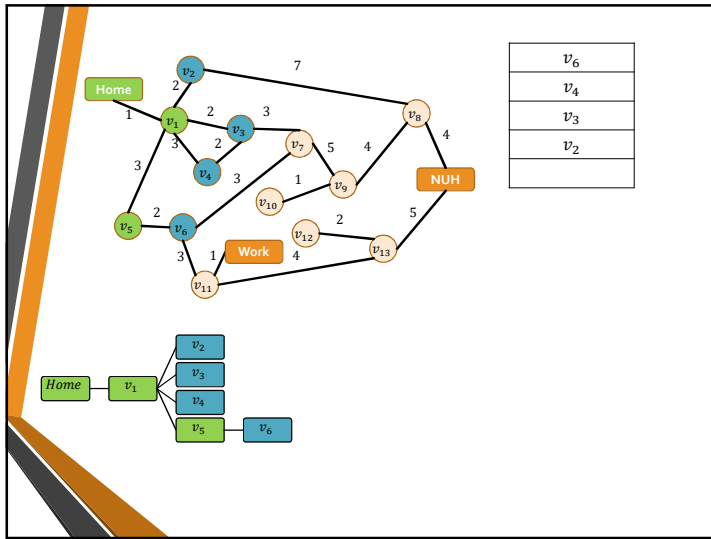
Uniform Cost Search

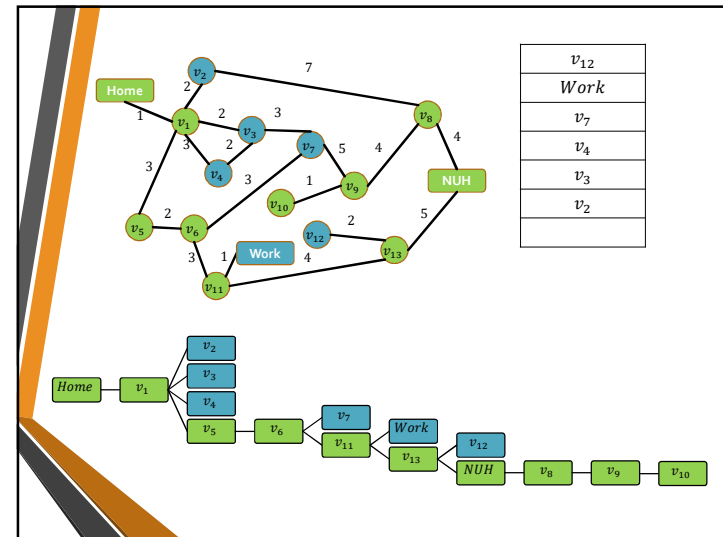
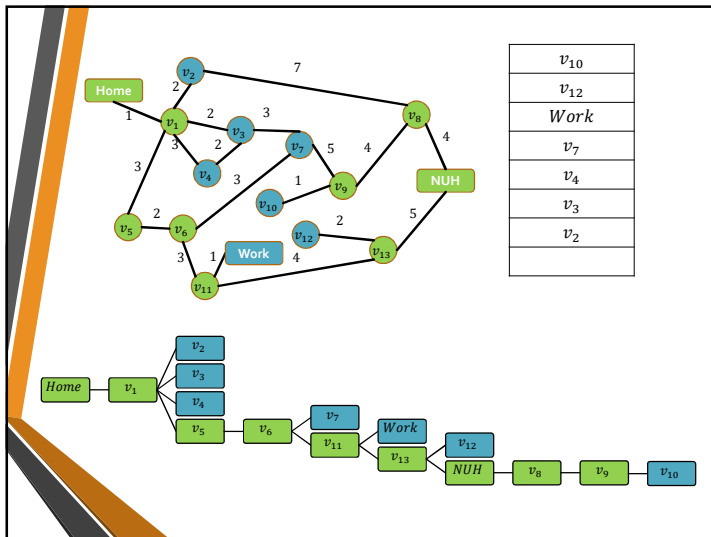
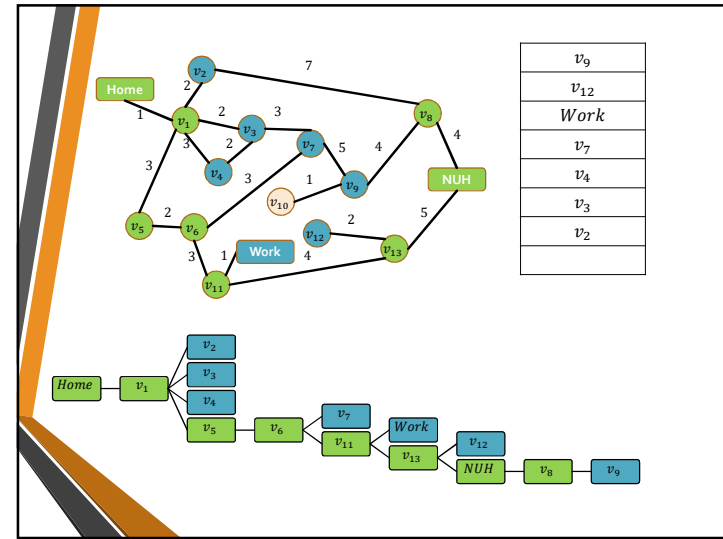
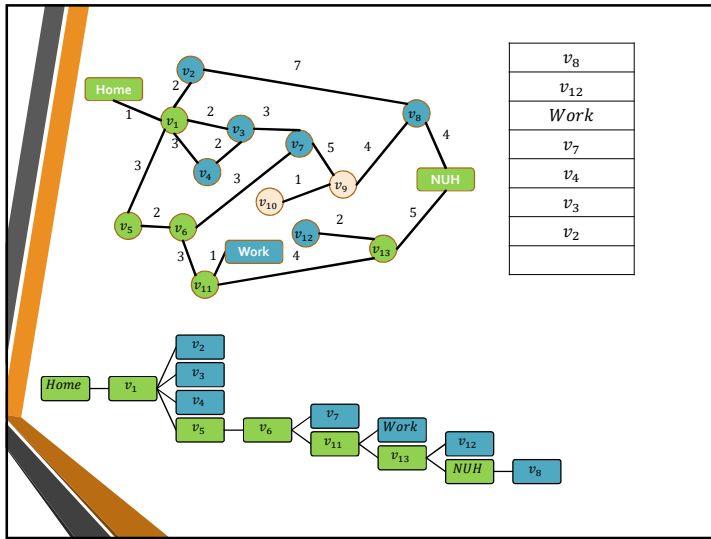
- At every round we get at least a distance of ϵ closer to the goal.
- Reach nodes at distance $0, \epsilon, 2\epsilon, \dots, \lceil \frac{C^*}{\epsilon} \rceil \epsilon$ of goal; total $\lceil \frac{C^*}{\epsilon} \rceil + 1$ steps.
- At step k (depth k at most): keep $\leq b^k$ nodes in frontier.

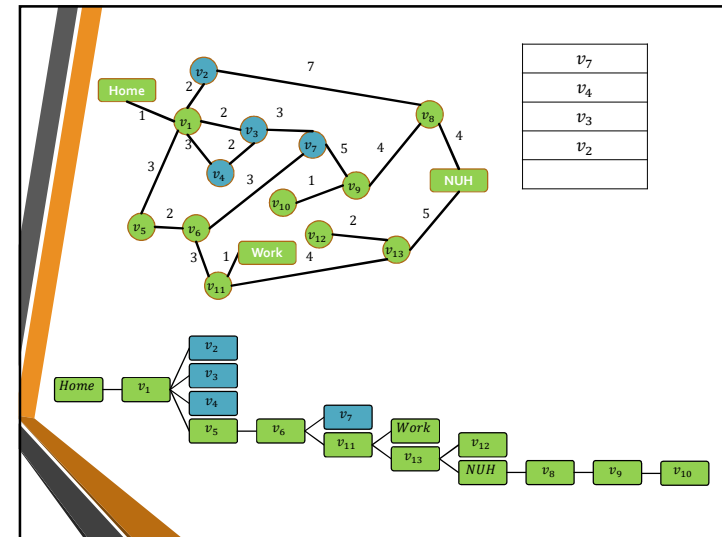
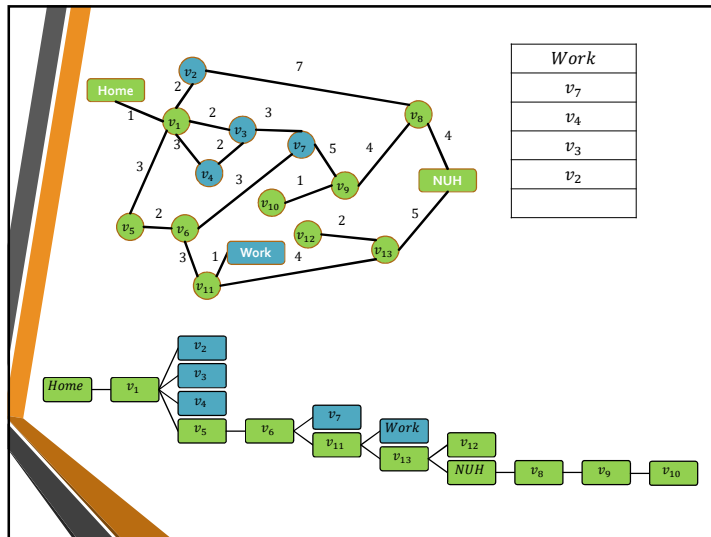
Depth-First Search

- Idea: Expand deepest unexpanded node
- Implementation: Frontier = LIFO stack, i.e., insert successors at the front









Properties of DFS

Property	
Complete?	No on infinite depth graphs
Optimal	No
Time	$\mathcal{O}(b^m)$
Space	$\mathcal{O}(bm)$ (can be $\mathcal{O}(m)$)

completeness: always find a solution if exists
 optimality: find a least-cost solution
 time complexity: number of nodes generated
 space complexity: max. number of nodes in memory

Depth-First Search

- When checking a node v , we push at most b descendants to stack.
- We do so at most m times $\Rightarrow \mathcal{O}(bm)$ space.
- Do we really need to push **all** b descendants to the stack?

Depth-Limited Search (DLS)

- Idea: run DFS with depth limit ℓ , i.e., do not search at depth greater than ℓ .
- Same guarantees as DFS, with ℓ instead of m ($\mathcal{O}(b^\ell)$ time; $\mathcal{O}(b^\ell)$ space)

Iterative Deepening Search (IDS)

- Idea: Perform DLSs with increasing depth limit until goal node is found
- Better if state space is large and depth of solution is unknown
- Implementation:

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
```

How Wasteful is IDS?

- Number of nodes generated in DLS to depth ℓ with branching factor b :

$$\mathcal{O}(b^0) + \mathcal{O}(b^1) + \mathcal{O}(b^2) + \dots + \mathcal{O}(b^{\ell-2}) + \mathcal{O}(b^{\ell-1}) + \mathcal{O}(b^\ell)$$

- Number of nodes generated in IDS to depth d with branching factor b :

$$(d+1)\mathcal{O}(b^0) + d\mathcal{O}(b^1) + (d-1)\mathcal{O}(b^2) + \dots + 3\mathcal{O}(b^{d-2}) + 2\mathcal{O}(b^{d-1}) + \mathcal{O}(b^d)$$

- For $b = 10, d = 5$,
 - $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
 - $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$
- Overhead = $\frac{123,456 - 111,111}{111,111} = 11\%$

Properties of IDS

Property	
Complete?	Yes (if b is finite)
Optimal	No (unless step cost is 1)
Time	$\mathcal{O}(b^d)$
Space	$\mathcal{O}(bd)$ (can be $\mathcal{O}(d)$)

completeness: always find a solution if exists
 optimality: find a least-cost solution
 time complexity: number of nodes generated
 space complexity: max. number of nodes in memory

All in

Property	BFS	UCS	DFS	DLS	IDS
Complete	Yes ¹	Yes ²	No	No	Yes ¹
Optimal	No ³	Yes	No	No	No ³
Time	$\mathcal{O}(b^d)$	$\mathcal{O}\left(b^{1+\left\lceil\frac{C^*}{\epsilon}\right\rceil}\right)$	$\mathcal{O}(b^m)$	$\mathcal{O}(b^\ell)$	$\mathcal{O}(b^d)$
Space	$\mathcal{O}(b^d)$	$\mathcal{O}\left(b^{1+\left\lceil\frac{C^*}{\epsilon}\right\rceil}\right)$	$\mathcal{O}(bm)$	$\mathcal{O}(b\ell)$	$\mathcal{O}(bd)$

1. BFS and IDS are complete if b is finite.
2. UCS is complete if b is finite and step cost $\geq \epsilon$
3. BFS and IDS are optimal if all step costs are identical

Choosing a Search Strategy

- Depends on the problem:
 - Finite/infinite depth of search tree
 - Known/unknown solution depth
 - Repeated states
 - Identical/non-identical step costs
 - Completeness and optimality needed?
 - Resource constraints (e.g., time, space)

83

Can We Do Better?

- Yes! Exploit problem-specific knowledge; obtain heuristics to guide search
- How?
 - Informed (heuristic) search
 - Expand “more promising” nodes.

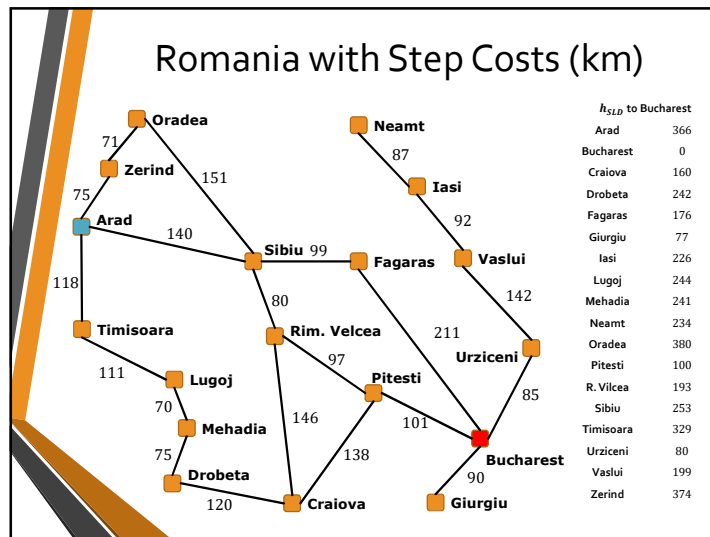
84

Best-First Search

- Idea: use an **evaluation function** $f(n)$ for each node n
 - Cost estimate \rightarrow Expand node with lowest evaluation/cost first
- Implementation:

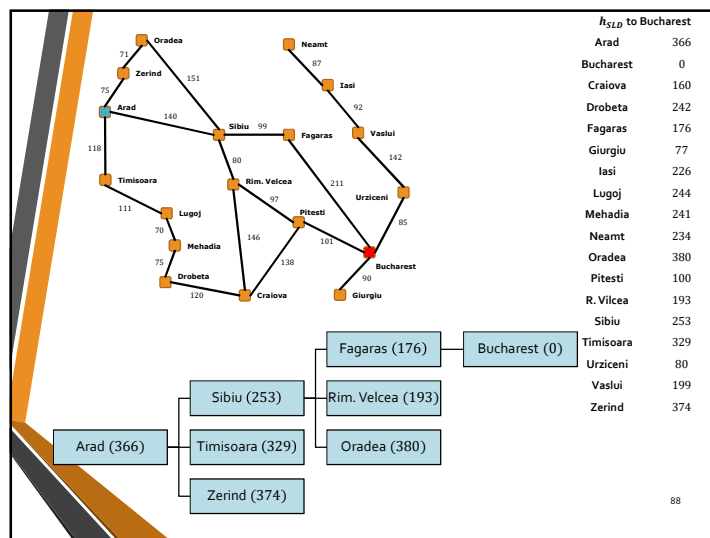
Frontier = priority queue ordered by nondecreasing cost f
- Special cases (different choices of f):
 - Greedy best-first search
 - A* search

85



Greedy Best-First Search

- Evaluation function
 $f(n) = h(n)$ (heuristic function) = estimated cost of cheapest path from n to goal
- e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal



Properties of Greedy Best-First Search

Property	
Complete?	Yes (if b is finite)
Optimal	No (shortest path to Bucharest: 418km)
Time	$\mathcal{O}(b^m)$, but a good heuristic can reduce complexity substantially
Space	Max size of frontier $\mathcal{O}(b^m)$

What Important Information Does the Algorithm Ignore?