

CS3243: Introduction to Artificial Intelligence

Semester 2, 2018/2019

Previously...

- **Uninformed** search strategies use only the information available in the problem definition
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search
- This class –exploit additional information!

Property	BFS	UCS	DFS	DLS	IDS
Complete	Yes ¹	Yes ²	No	No	Yes ¹
Optimal	No ³	Yes	No	No	No ³
Time	$\mathcal{O}(b^d)$	$\mathcal{O}\left(b^{1+\lceil\frac{C^*}{\epsilon}\rceil}\right)$	$\mathcal{O}(b^m)$	$\mathcal{O}(b^\ell)$	$\mathcal{O}(b^d)$
Space	$\mathcal{O}(b^d)$	$\mathcal{O}\left(b^{1+\lceil\frac{C^*}{\epsilon}\rceil}\right)$	$\mathcal{O}(bm)$	$\mathcal{O}(b\ell)$	$\mathcal{O}(bd)$

1. BFS and IDS are complete if b is finite.
2. UCS is complete if b is finite and step cost $\geq \epsilon$
3. BFS and IDS are optimal if step costs are identical.

Q1. Consider the following parameters
 $b = 3, d = 3$, what is the overhead of
 IDS as compared to DLS?

Q2. Consider the following parameters
 $b = 3, d = 3$, what is the number of
nodes to keep track in DFS?

Choosing a Search Strategy

- Depends on the problem:
 - Finite/infinite depth of search tree
 - Known/unknown solution depth
 - Repeated states
 - Identical/non-identical step costs
 - Completeness and optimality needed?
 - Resource constraints (e.g., time, space)

Can We Do Better?

- Yes! Exploit problem-specific knowledge; obtain heuristics to guide search
- Today:
 - Informed (heuristic) search
 - Expand “more promising” nodes.

INFORMED SEARCH

AIMA Chapter 3.5.1–3.5.2, 3.6.1–3.6.2


Outline

- Best-first search
- Greedy best-first search
- A* search
- Heuristics

Best-First Search

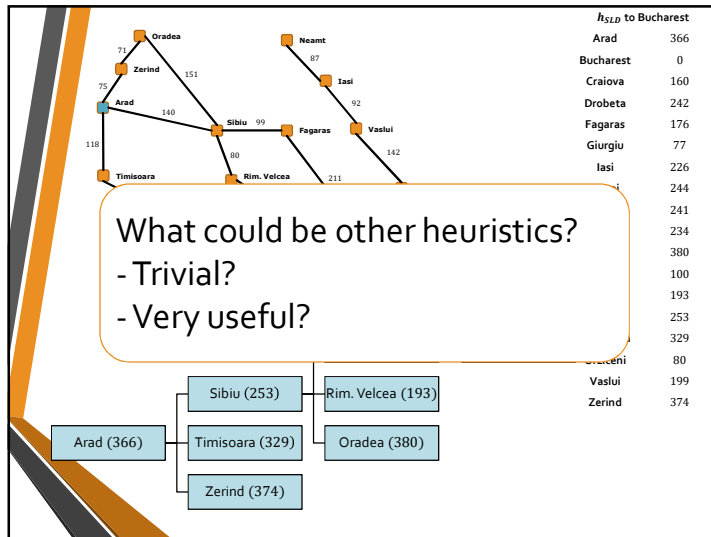
- Idea: use an **evaluation function** $f(n)$ for each node n
 - Cost estimate \rightarrow Expand node with lowest evaluation/cost first
- Implementation:
 - Frontier = priority queue ordered by nondecreasing cost f
- Special cases (different choices of f):
 - Greedy best-first search
 - A* search

Romania with Step Costs (km)



Greedy Best-First Search

- Evaluation function $f(n) = h(n)$ (**h**euristic function) = estimated cost of cheapest path from n to goal
 - Restriction: $h(goal) = 0$
- e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal



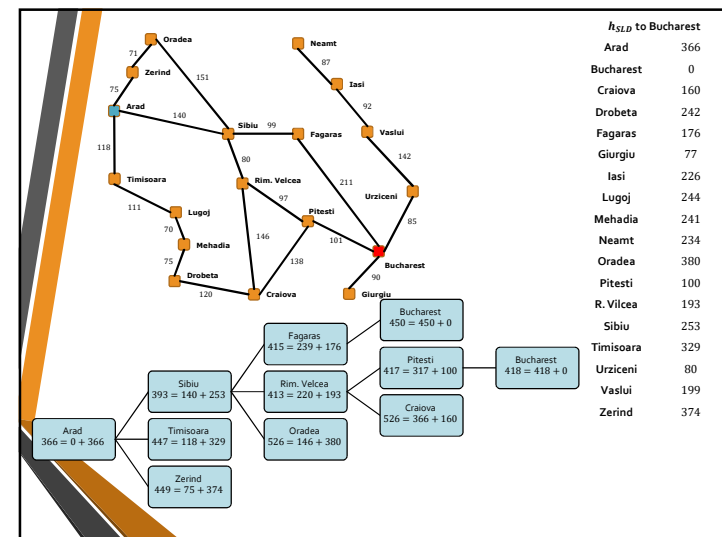
Properties of Greedy Best-First Search

Property	
Complete?	Yes (if b is finite)
Optimal	No (shortest path to Bucharest: 418km)
Time	$\mathcal{O}(b^m)$, but a good heuristic can reduce complexity substantially
Space	Max size of frontier $\mathcal{O}(b^m)$

What Important Information Does the Algorithm Ignore?

A* Search

- Idea: Avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
- $g(n)$ = cost of reaching n from start node
- $h(n)$ = cost estimate from n to goal
- $f(n)$ = estimated cost of cheapest path **through** n to goal



Admissible Heuristics

- $h(n)$ is **admissible** if, $\forall n, h(n) \leq h^*(n)$
- $h^*(n)$ = **true** cost to reach the goal state from n .
- **Never overestimates** cost to reach goal
- Example: $h_{SLD}(n)$ never overestimates the actual road distance (roads are at best straight!)

Admissible Heuristics

Theorem: If $h(n)$ is admissible, then A^* using TREE-SEARCH is optimal


Optimality of A^* using TREE-SEARCH

t - a suboptimal goal in the frontier.

n - an unexpanded node in the frontier; n is on a shortest path to an optimal goal t^* .

It would be **very** bad if suboptimal goal node t gets checked before n !!

$\Rightarrow f(t)$ is lower than $f(n)$




Optimality of A^* using TREE-SEARCH

t - a suboptimal goal in the frontier.

n - an unexpanded node in the frontier; n is on a shortest path to an optimal goal t^* .

t gets checked after n if $f(t) > f(n)$

We need to show that t gets checked after n



Optimality of A* using TREE-SEARCH

t - a suboptimal goal in the frontier.


n - an unexpanded node in the frontier; n is on a shortest path to an optimal goal t^* .

t gets checked after n if
 $g(t) + h(t) > g(n) + h(n)$

In A*, $f(v) = g(v) + h(v)$

Cost to get
to v from s

Est. dist. from
 v to goal




Optimality of A* using TREE-SEARCH

t - a suboptimal goal in the frontier.

n - an unexpanded node in the frontier; n is on a shortest path to an optimal goal t^* .

t gets checked after n if
 $g(t) + h(t) > g(n) + h(n)$

This is 0




Optimality of A* using TREE-SEARCH

t - a suboptimal goal in the frontier.

n - an unexpanded node in the frontier; n is on a shortest path to an optimal goal t^* .

t gets checked after n if
 $g(t) > g(n) + h(n)$

$$\begin{aligned} f(t) &= g(t) > g(t^*) \\ &= g(n) + d(n, t^*) \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$




Consistent Heuristics

- A heuristic is **consistent** if, for every node n and every successor n' of n generated by any action a ,

$$h(n) \leq c(n, n') + h(n')$$

Lemma: if h is consistent,

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$




$f(n)$ is non-decreasing along any path.

Consistent Heuristics

Theorem: If $h(n)$ is consistent, then A^* using GRAPH-SEARCH is optimal

Optimality of A^* using GRAPH-SEARCH


- $f(n)$ is non-decreasing along any path
- A^* expands nodes in non-decreasing order of f value
- Gradually adds " f -contours" of nodes
- Contour i has all nodes with $f = f_i$ where $f_i < f_{i+1}$



Optimality of A^* using GRAPH-SEARCH

Stronger claim: when A^* selects a node n for expansion, **the shortest path to n has been found** (proof by induction on distance from s).

Assume otherwise; let n be the first node reached by suboptimal path \Rightarrow all nodes before n reached by optimal path.



Properties of A^* Search


Property	
Complete?	Yes (if there is a finite no. of nodes with $f(n) \leq f(G)$)
Optimal	Yes
Time	$\mathcal{O}(b^{h^*(s_0)-h(s_0)})$ where $h^*(s_0)$ is actual cost of getting from root to goal.
Space	Max size of frontier $\mathcal{O}(b^m)$

Admissible vs. Consistent Heuristics

- Why is consistency a stronger sufficient condition than admissibility?
 - Consistent \Rightarrow admissible
 - Admissible \nRightarrow consistent
- $k(n)$ be the cost of cheapest path from n to goal
- To prove, $h(n) \leq k(n)$
 - If n is the goal, $h(n) = 0 \leq k(n)$
 - Induction hypothesis: $h(n') \leq k(n')$
 - Let, path from n to goal has i steps, n' is the successor of n
 - $n' \rightsquigarrow$ goal, $i - 1$ steps
- h is consistent \Rightarrow we have: $h(n) \leq c(n, n') + h(n')$
 - $\leq c(n, n') + k(n') = k(n)$


Admissible vs. Consistent Heuristics

- An admissible but inconsistent heuristic cannot guarantee optimality of A* using GRAPH-SEARCH
 - GRAPH-SEARCH discards new paths to a repeated state. May discard the optimal path.
 - Consistent heuristic: always follows optimal path (that lemma was important!)



Admissible Heuristics


- Let's revisit the 8-puzzle
 - Branching factor is about 3
 - Average solution depth is about 22 steps
 - Exhaustive tree search examines 3^{22} states
- How do we come up with good heuristics?



Admissible Heuristics

E.g., 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance (i.e., no. of squares from desired location of each tile)



$h_1(s) = 8$

$h_2(s) = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18$

Dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible), then h_2 **dominates** h_1 . It follows that h_2 incurs lower search cost than h_1 .

Average search costs (nodes generated):

$d = 12$

Algorithm	# Nodes
IDS	3,644,035
$A^*(h_1)$	227
$A^*(h_2)$	73

$d = 24$

Algorithm	# Nodes
IDS	Galactic Number
$A^*(h_1)$	39,135
$A^*(h_2)$	1,641

Deriving Admissible Heuristics

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem

Deriving Admissible Heuristics

- Rules of 8-puzzle:

A tile can move from square A to square B if A is horizontally or vertically adjacent to B and B is blank

- We can generate three relaxed problems

- A tile can move from square A to square B if A is adjacent to B
- A tile can move from square A to square B if B is blank
- A tile can move from square A to square B

Deriving Admissible Heuristics

- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ is the resulting heuristic.
- If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ (Manhattan Dist.) is the resulting heuristic

LOCAL SEARCH


AIMA Chapter 4.1

Local Search Algorithms

- The **path** to goal is irrelevant; the goal state itself is the solution
- State space = set of “complete” configurations
- Find final configuration satisfying constraints, e.g., n -queens
- **Local search algorithms**: maintain single “current best” state and try to improve it
- Advantages:
 - very little/constant memory
 - find reasonable solutions in large state space

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal




Hill-Climbing Search

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  current ← MAKE-NODE(problem.INITIAL-STATE)
  loop do
    neighbor ← a highest-valued successor of current
    if neighbor.VALUE ≤ current.VALUE then return current.STATE
    current ← neighbor
```

“Like climbing Mt. Everest in thick fog with amnesia”

Hill-Climbing Search

- Problem: depending on initial state, can get stuck in local maxima




- Non-guaranteed fixes: sideways moves, random restarts

Hill-Climbing Search: 8-Queens

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$ for the above state

Hill-Climbing Search: 8-Queens



Local Minimum with $h = 1$

Hill-climbing

- Use hill-climbing if we are OK with approximate solutions
- If accuracy matters, need to be more careful



Local search strategies

- Hill-climbing search: use of heuristic function to improve “current” state