

LOCAL SEARCH

AIMA Chapter 4.1

Previously...


Use Heuristics to Guide Search

- Greedy best-first search
- A^* search

A^* Search Heuristic


- If $h(n)$ is admissible, then A^* w. Tree-Search is optimal
- If $h(n)$ is consistent, then A^* w. Graph-Search is optimal
- A heuristic that dominates another incurs lower search cost

Consider the search problem



	$h(A)$	$h(B)$	$h(G)$
H-1	4	3	0
H-2	5	2	0

Consider the search problem



Consistent?

$$h(A) - h(B) \leq \text{pathCost}$$

Admissible?

$$h(A) \leq \text{pathToGoal}^*$$


	$h(A)$	$h(B)$	$h(G)$
H-1	4	3	0
H-2	5	2	0

Local Search Algorithms

- The **path** to goal is irrelevant; the goal state itself is the solution
- State space = set of “complete” configurations
- Find final configuration satisfying constraints, e.g., n -queens
- **Local search algorithms**: maintain single “current best” state and try to improve it
- Advantages:
 - very little/constant memory
 - find reasonable solutions in large state space

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal




Hill-Climbing Search

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  current ← MAKE-NODE(problem.INITIAL-STATE)
  loop do
    neighbor ← a highest-valued successor of current
    if neighbor.VALUE ≤ current.VALUE then return current.STATE
    current ← neighbor
```

“Like climbing Mt. Everest in thick fog with amnesia”

Hill-Climbing Search

- Problem: depending on initial state, can get stuck in local maxima




- Non-guaranteed fixes: sideways moves, random restarts

Hill-Climbing Search: 8-Queens

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♙	13	16	13	16
♙	14	17	15	♙	14	16	16
17	♙	16	18	15	♙	15	♙
18	14	♙	15	15	14	♙	16
14	14	13	17	12	14	12	18

- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$ for the above state

Hill-Climbing Search: 8-Queens




Local Minimum with $h = 1$


Adversarial Search

AIMA Chapter 5.1 – 5.5

AI vs. Human Players: the State of the Art



Kasparov Vs. IBM's Deep Blue (1997)




World Poker Champs Vs. Libratus (2017)



Lee Sedol Vs. DeepMind's AlphaGo (2016)

AI vs. Human Players: the State of the Art



World Poker Champs Vs. Libratus (2017)



Lee Sedol Vs. DeepMind's AlphaGo (2016)



AlphaZero can solve any zero-sum game! (2018)


AI vs. Human Players: the State of the Art



Lee Sedol Vs. DeepMind's AlphaGo (2016)



AlphaZero can solve any zero-sum game! (2018)



OpenAI Five
OpenAI 5 vs world champs—DOTA2 (2019)

Deterministic Games in Practice

- [Chinook \(Checkers, 1994\)](#). Precomputed endgame database \Rightarrow perfect play for all positions with ≤ 8 pieces on the board (total of 444 billion positions).
- [Deep Blue \(Chess, 1997\)](#). Searches 200 million positions/sec + evaluation functions + secret sauce.
- [Logistello \(Othello, 1997\)](#). Human champions refuse to play against AI.
- [AlphaGo + Alphazero \(Go/everything above, 2016-2017\)](#). Learning for evaluation functions + database and efficient search + secret sauce.

Outline

- Adversarial search problems (aka games)
- Optimal (i.e., Minimax) decisions
- α - β pruning
- Imperfect, real-time decisions

Games vs. Search Problems

Utility maximizing opponent

- solution is a strategy specifying a move for every possible opponent response.

Time limit

- unlikely to find goal, must approximate

Let's Play!

- Two players in a **zero-sum game**:
 - Winner gets paid and loser pays.
- Easy to think in terms of a **max player** and **min player**
 - Player 1 wants to maximize value (MAX player)
 - Player 2 wants to minimize value (MIN player)



Game: Problem Formulation

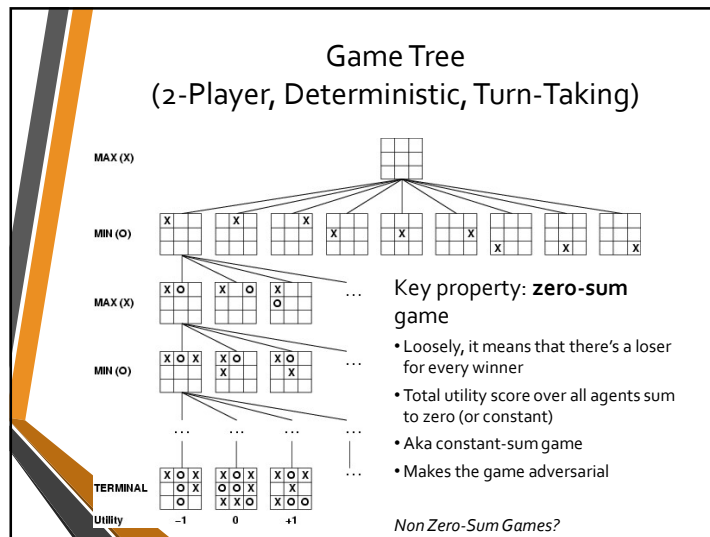
A game is defined by 7 components:

1. Initial state s_0
2. States
3. Players: $PLAYER(s)$ defines which player has the move in state s .
4. Actions : $ACTIONS(s)$ returns set of legal moves in state s
5. Transition model : $RESULT(s, a)$ returns state that results from the move a in state s .

Game: Problem Formulation

A game is defined by 7 components:

6. Terminal test $TERMINAL(s)$ returns true if game is over and false otherwise
 - Terminal states: states where the game has ended
7. Utility function $UTILITY(s, p)$ gives final numeric value for a game that ends in terminal state s for a player p
 - Chess: White wins +1; Black wins -1; draw 0.



Example : Game of NIM

Several piles of sticks are given. We represent the configuration of the piles by a monotone sequence of integers, such as (1,3,5). A player may remove, in one turn, any number of sticks from one pile. Thus, (1,3,5) would become (1,1,3) if the player were to remove 4 sticks from the last pile. The player who takes the last stick loses.


Example : Game of NIM

Start	(A)	(B)	(C)
Bob take 2 from A			
Alice take 3 from C			
Bob take 1 from B			
Alice take 1 from B			
Bob take 1 from A	o		
Alice take 1 from B	o		
Bob take 2 from C	o		o

Bob wins!!

||| | |||

Represent the NIM game (1,2,2) as a game tree.



Represent the NIM game $(1,2,2)$ as a game tree.

Play against what kind of player?
Stupid? Smart?

Player Strategies

A strategy s for player i : what will player i do at every node of the tree that they make a move in?

Need to specify behavior in states that will never be reached!

Winning Strategy

A strategy s_1^* for player 1 is called **winning** if for any strategy s_2 by player 2, the game ends with player 1 as the winner.

A strategy t_1^* for player 1 is called **non-losing** if for any strategy s_2 by player 2, the game ends in either a tie or a win for player 1.

Nash Equilibrium

When players know the strategies of all opponents: no one wants to change her strategy

Stronger form of Nash equilibrium – subgame perfect Nash equilibrium

Every subgame is a Nash equilibrium


Optimal Strategy at Node - Minimax

$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s) & \text{if TerminalTest}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{MIN} \end{cases}$$


Intuitively,

- MAX chooses move to maximize the minimum payoff
- MIN chooses move to minimize the maximum payoff

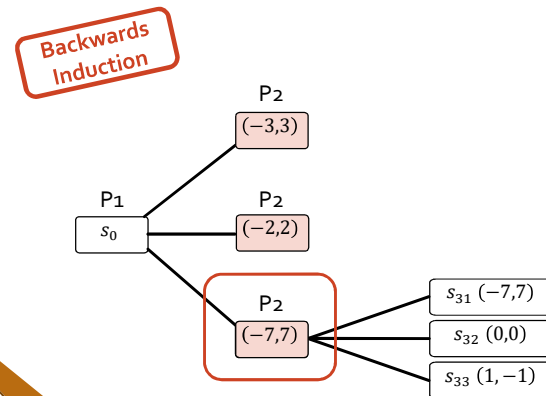
Minimax Play (Subperfect Nash Equilibrium)

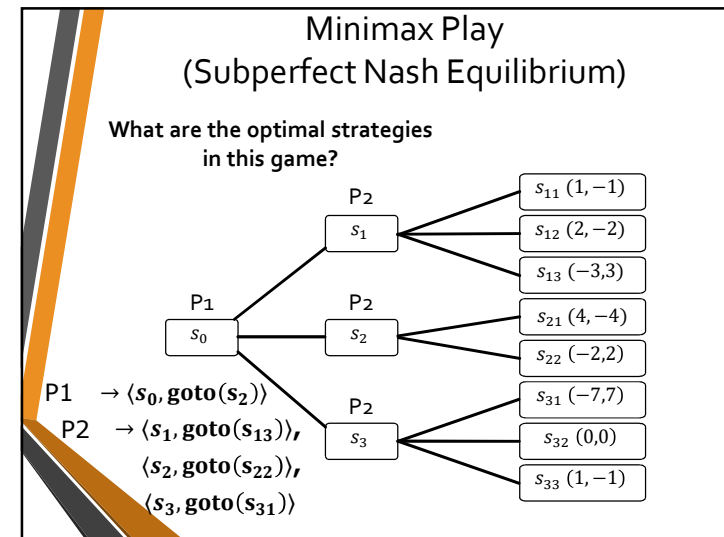
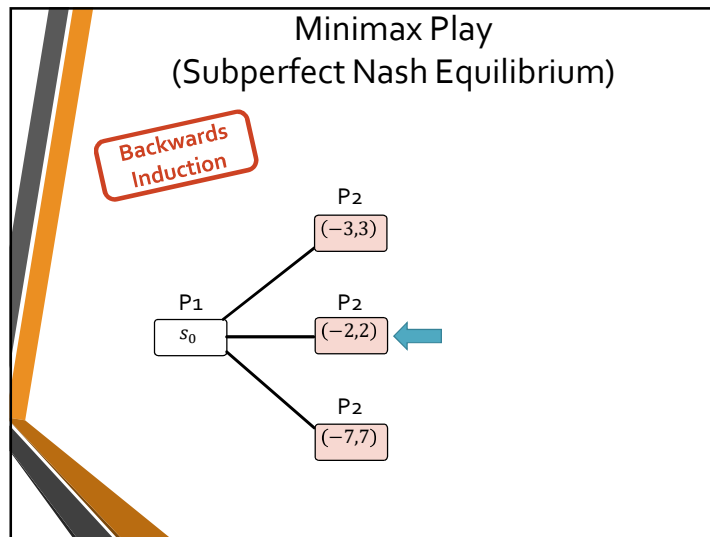


Minimax Play (Subperfect Nash Equilibrium)



Minimax Play (Subperfect Nash Equilibrium)





Properties of Minimax

Property	Computes a Nash equilibrium of the game
Complete?	Yes (if game tree is finite)
Optimal	Yes (optimal gameplay)
Time	$\mathcal{O}(b^m)$
Space	Like DFS: $\mathcal{O}(bm)$

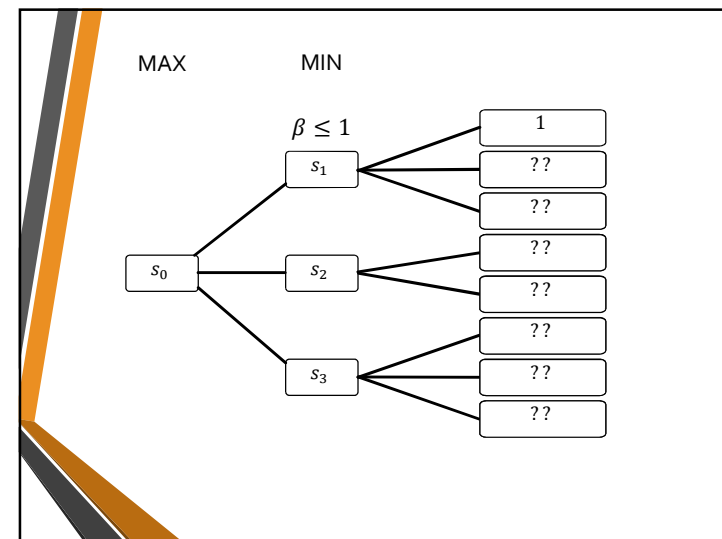
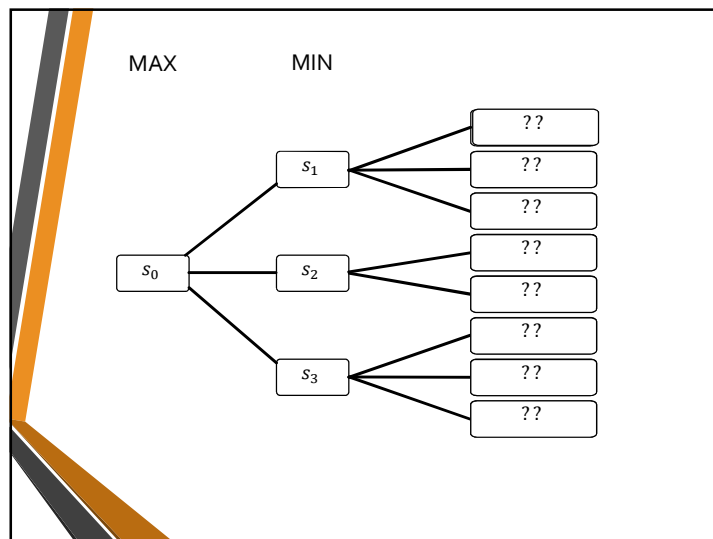
- ### Minimax Algorithm
- Runs in time polynomial in tree size
 - Returns a sub-perfect Nash equilibrium: the best action at every choice node.
- Are we done here?**

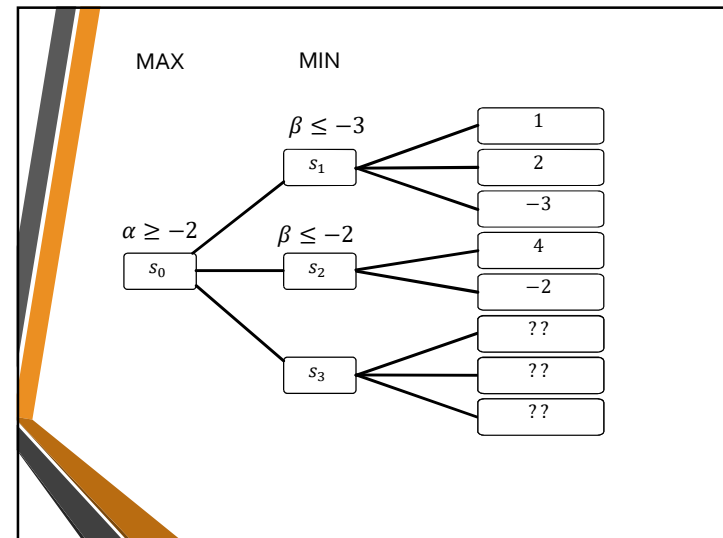
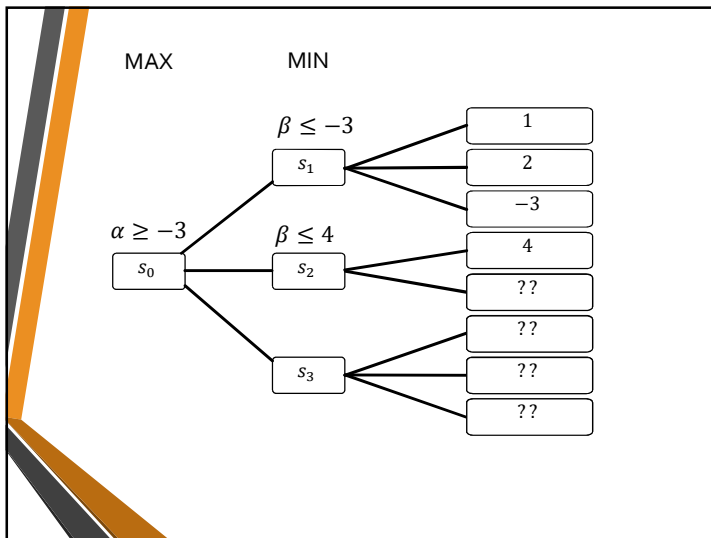
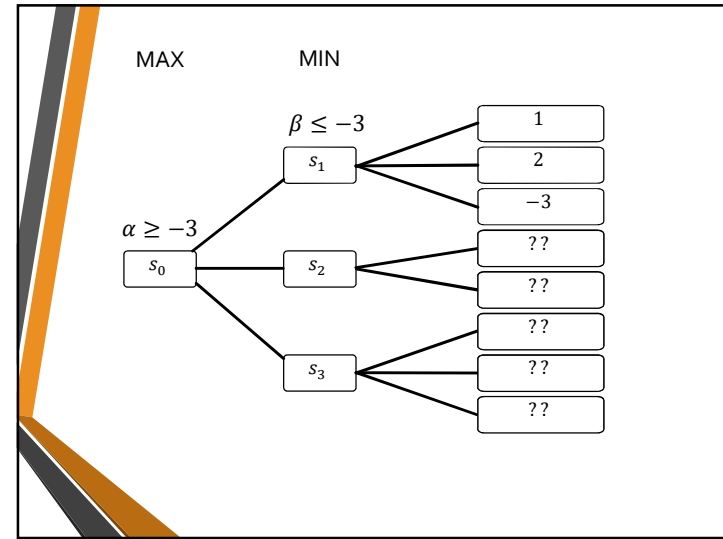
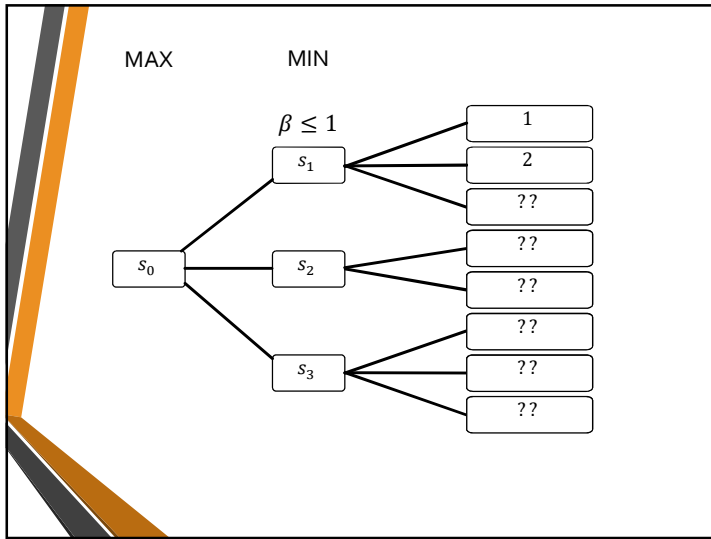
Backwards Induction

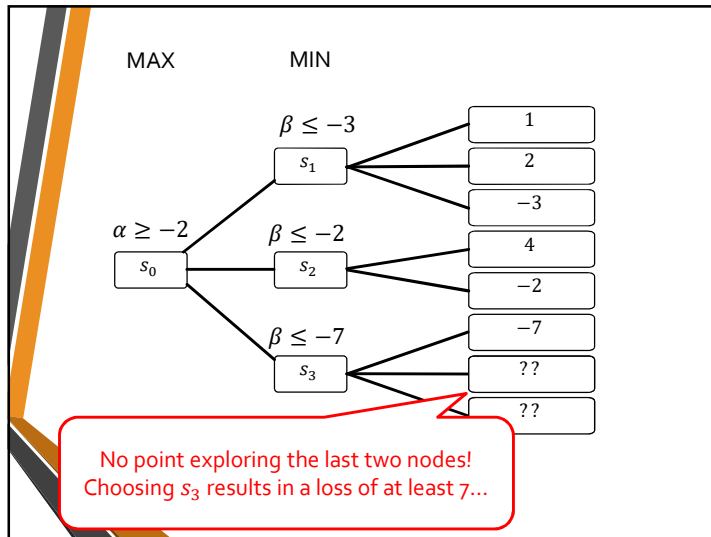
- Game trees are huge: chess has game tree with $\sim 10^{123}$ nodes (planet Earth has $\sim 10^{50}$ atoms)
- Impossible to expand the entire tree

α - β Pruning

- **Basic idea:** "If you have an idea that is surely bad, don't take the time to see how truly awful it is." -- Pat Winston
- Maintain a **lower bound α** and **upper bound β** of the values of, respectively, MAX's and MIN's nodes seen thus far. We can prune subtrees that will never affect minimax decision.







α - β Pruning

- MAX node n : $\alpha(n)$ = highest **observed** value found on path from n ; initially $\alpha(n) = -\infty$
- MIN node n : $\beta(n)$ is the lowest observed value found on path from n ; initially $\beta(n) = +\infty$
- Pruning:**
 - given a MIN node n , stop searching below n if there is some MAX ancestor i of n with $\alpha(i) \geq \beta(n)$
 - given a MAX node n , stop searching below n if there is some MIN ancestor i of n with $\beta(i) \leq \alpha(n)$

Analysis of α - β Pruning

- When we prune a branch, it **never** affects final outcome.
- Good move ordering improves effectiveness of pruning
- "Perfect" ordering: time complexity = $O(b^{\frac{m}{2}})$
 - Good pruning strategies allow us to search twice as deep!
 - Chess: simple ordering (checks, then take pieces, then forward moves, then backwards moves) gets you close to best-case result.
 - It makes sense to have good expansion order heuristics.
- Random ordering: time complexity = $O(b^{\frac{3m}{4}})$ for $b < 1000$

Summary: α - β Pruning Algorithm

- Initially, $\alpha(n) = -\infty$, $\beta(n) = +\infty$
- $\alpha(n)$ is max along search path containing n
- $\beta(n)$ is min along search path containing n
- If a MIN node has value $v \leq \alpha(n)$, no need to explore further.
- If a MAX node has value $v \geq \beta(n)$, no need to explore further.

Time Limit

- **Problem:** very large search space in typical games
- **Solution:** α - β pruning removes large parts of search space
- Unresolved problem: Maximum depth of tree
- Standard solutions:
 - **evaluation function** = estimated expected utility of state
 - **cutoff test:** e.g., depth limit

Heuristic Minimax Value

$\text{MINIMAX}(s) =$
 $\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$

$\text{H-MINIMAX}(s, d) =$
 $\begin{cases} \text{EVAL}(s) & \text{if } \text{CUTOFF-TEST}(s, d) \\ \max_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$

Run minimax until depth d ; then start using the evaluation function to choose nodes.

Evaluation Functions

- An evaluation function is a mapping from game states to real values: $f: \mathcal{S} \rightarrow \mathbb{R}$
 - So far:

$$f(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL}(s) \\ 0 & \text{otherwise} \end{cases}$$
- Should be cheap to compute
- For non-terminal states, must be strongly correlated with actual chances of winning
- Tic-Tac-Toe:

$$f(n) = [\text{\# of open lines for X}] - [\text{\# of open lines for O}]$$

Evaluation Functions

- Chess:
 - Alan Turing's evaluation function:

$$f(n) = \frac{w(n)}{b(n)}$$
 where $w(n)$ is the point value of white's pieces and $b(n)$ is the point value of black's pieces.
 - Modern evaluation functions: weighted sum of position features

$$f(n) = w_1 F_1(n) + w_2 F_2(n) + \dots + w_k F_k(n)$$
 - Example features: piece count, piece placement, controlled squares etc.
 - Deep Blue has about 6000 features.
 - How do we determine weights? Do they change dynamically?

Evaluation Functions

- Suppose that $f(n) = \sum_{j=1}^k w_j F_j(n)$
- If for $n, n' \in \mathcal{S}$ we have $F_j(n) = F_j(n')$ for all j , then they're indistinguishable.
- ... evaluates all states with same value.

- Let $\mathcal{S}(\vec{q}) = \{n \in \mathcal{S} : \forall j \in [k], F_j(n) = q_j\}$: all states whose feature values are as specified in the vector \vec{q} .
- "All states where white has two pawns but black has a bishop."
- Suppose we know that in this case,
 - Black wins 61% of games
 - White wins 15% of games
 - 24% of games end in a draw
- Then expected utility for black is
$$0.61 \times 1 + 0.15 \times (-1) + 0.24 \times 0 = 0.46$$
- Evaluation function need not return actual expected values, just maintain relative order of states.

Evaluation alone is not enough!

- Opening move search doesn't result in useful utility estimates
- Applying evaluation functions on end game scenarios may not **solve** the game
 - Western chess KRK end game
- Use a **policy** or lookup table (taken from expert knowledge or previous game history)

Cutting Off Search

- Modify minimax or α - β pruning algorithms by replacing
 - `TERMINAL-TEST(state)` with `CUTOFF-TEST(state, depth)`
 - `UTILITY(state)` is replaced by `EVAL(state)`
- Can also be combined with iterative deepening


Cutting Off Search

- Why is it useful?
 - For chess, $b^m = 10^6$, $b = 35 \Rightarrow m = 4$
 - Ply = 4 – Novice chess player
 - Ply = 8 – Typical PC engine
 - Ply = 12~14 – Grandmaster chess (Deep blue)

Stochastic Games

- Many games have randomization:
 - Backgammon
 - Settlers of Catan
 - Poker
- How do we deal with uncertainty?
- Can we still use minimax? Yes, but search space is much bigger

Adding Chance Layers



Calculate the **expected** value of a state
(MUCH harder than deterministic games)