

CS2105

An **Awesome** Introduction to Computer Networks


Lectures 2&3: The Application Layer



© CS2105


Packet Delay

- ❖ End-to-end delay is the time taken for a packet to travel from source to destination. It consists of:
 - processing delay
 - queueing delay
 - transmission delay
 - propagation delay



© CS2105

Lectures 2&3 - 5



© CS2105

Lectures 2&3 - 7

Lectures 2&3: Roadmap

2.1 Principles of Network Applications

2.2 Web and HTTP

2.4 DNS

2.7 Socket programming

To discuss next week


Kurose Textbook, Chapter 2
(Some slides are taken from the book)

© CS2105

Packet Switching

- ❖ The Internet is a **packet switching** network
 - Hosts share and contend network resources.
 - **Application message** is broken into a bunch of packets and sent onto the link one by one.
 - A router stores and forwards packets.
 - Receiver assembles all the packets to restore the **application message**.

Bandwidth division into "pieces"
Dedicated allocation
Resource reservation



© CS2105

Network Protocols

- ❖ Networks are complex. There are many issues to consider, to support different applications running on large number of hosts through different access technology and physical media.
- ❖ **Protocols** regulate communication activities in a network.
 - Define the **format** and **order** of messages exchanged between hosts for a specific purpose.

© CS2105

Lectures 2&3: Application Layer

After this class, you are expected to:

- ❖ understand the basic HTTP interactions between the client and the server, the concepts of persistent and non-persistent connections.
- ❖ understand the services provided by DNS and how a DNS query is resolved.
- ❖ understand the concept of socket.
- ❖ be able to write simple client/server programs through socket programming.

© CS2105

Evolution of Network Applications

- ❖ Early days of Internet
 - Remote access (e.g. telnet, now ssh)
- ❖ 70s – 80s
 - Email, FTP
- ❖ 90s
 - Web
- ❖ 2000s – now
 - P2P file sharing
 - Online games
 - Instant Messaging, Skype
 - YouTube, Facebook

Lectures 2&3 - 6

Lectures 2&3 - 8

PREVIOUS LECTURE

PREVIOUS LECTURE

PREVIOUS LECTURE


Creating Network Applications

write programs that:

- ❖ run on (different) *hosts*
- ❖ communicate over network
- ❖ e.g., web server software ↔ browser software

classic structure of network applications:

- ❖ Client-server
- ❖ Peer-to-peer (P2P)




© CS2105

Lectures 2&3 - 11

P2P Architecture

- ❖ No always-on server
- ❖ Arbitrary end systems directly communicate.
- ❖ Peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- ❖ Peers are intermittently connected and change IP addresses
 - complex management



© CS2105

Lectures 2&3 - 13

Requirements of Example Apps

Application	Data loss	Throughput	Time-sensitive
File transfer	No loss	Elastic	No
Electronic mail	No loss	Elastic	No
Web documents	No loss	Elastic	No
Real-time audio/video	Loss-tolerant	Audio: 5kbps-1Mbps Video: 10kbps-5Mbps	Yes: 100s of msec
Stored audio/video	Loss-tolerant	Same as above	Yes: few seconds
Interactive games	Loss-tolerant	Few kbps – 10 kbps	Yes: 100s of msec
Text messaging	No loss	Elastic	Yes and no

© CS2105

Lectures 2&3 - 15

Internet Transport Protocols

TCP service:

- ❖ *reliable transport* between sending and receiving process
- ❖ *flow control*: sender won't overwhelm receiver
- ❖ *congestion control*: throttle sender when network is overloaded
- ❖ *does not provide*: timing, minimum throughput guarantee, security

UDP service:

- ❖ *unreliable data transfer* between sending and receiving process
- ❖ *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee or security


Client-Server Architecture

Server:

- ❖ Waits for incoming requests
- ❖ Provides requested service to client
- ❖ data centers for scaling

Client:

- ❖ Initiates contact with server ("speaks first")
- ❖ Typically requests service from server
- ❖ For Web, client is usually implemented in browser



© CS2105

Lectures 2&3 - 12

What transport service does an app need?

Data integrity

- ❖ some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- ❖ other apps (e.g., audio streaming) can tolerate some data loss

Throughput

- ❖ some apps (e.g., multimedia) require minimum amount of bandwidth to be "effective"
- ❖ other apps (e.g., file transfer) make use of whatever throughput available

Timing

- ❖ some apps (e.g., online interactive games) require low delay to be "effective"

Security

- ❖ encryption, data integrity, authentication ...

© CS2105

Lectures 2&3 - 14

App-layer Protocols Define...

- ❖ *types of messages exchanged*
 - e.g., request, response
- ❖ *message syntax*:
 - what fields in messages & how fields are delineated
- ❖ *message semantics*
 - meaning of information in fields
- ❖ *rules* for when and how applications send & respond to messages
- ❖ *open protocols*:
 - defined in RFCs
 - allows for interoperability
 - e.g., HTTP, SMTP
- ❖ *proprietary protocols*:
 - e.g., Skype

© CS2105

Lectures 2&3 - 16

Example App/Transport Protocols

Application	Application Layer Protocol	Underlying Transport Protocol
Electronic mail	SMTP [RFC 5321]	TCP
Remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
File transfer	FTP [RFC 959]	TCP
Streaming multimedia	HTTP (e.g., YouTube)	TCP
Internet telephony	SIP [RFC 3261], RTP [RFC 3550], or proprietary (e.g., Skype)	TCP or UDP

Lectures 2&3: Roadmap

2.1 Principles of Network Applications

2.2 Web and HTTP

2.4 DNS

2.7 Socket programming

The Web: Some Jargon

- ❖ A Web page typically consists of:
 - *base HTML file*, and
 - *several referenced objects*.
- ❖ An object can be HTML file, JPEG image, Java applet, audio file, ...
- ❖ Each object is addressable by a *URL*, e.g.,

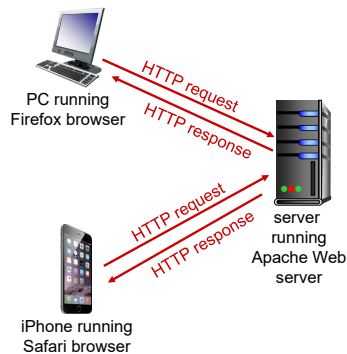
www.comp.nus.edu.sg/~cs2105/img/doge.jpg

host name path name

HTTP Overview

HTTP: Hypertext transfer protocol

- ❖ Web's application layer protocol
- ❖ Client/server model
 - *client*: usually is browser that requests, receives and displays Web objects
 - *server*: Web server sends objects in response to requests
- ❖ http 1.0: RFC 1945
- ❖ http 1.1: RFC 2616



HTTP Over TCP

HTTP uses TCP as transport service

- ❖ Client initiates TCP connection to server.
- ❖ Server accepts TCP connection request from client.
- ❖ HTTP messages are exchanged between browser (HTTP client) and Web server (HTTP server) over TCP connection.
- ❖ TCP connection closed.

Two Types of HTTP Connections

Non-persistent HTTP

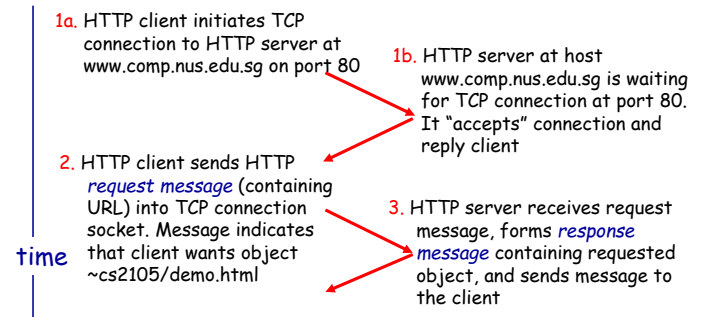
- ❖ at most one object sent over a TCP connection
 - connection then closed
- ❖ downloading multiple objects required multiple connections

Persistent HTTP

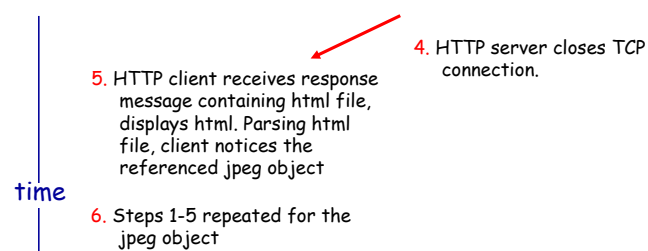
- ❖ multiple objects can be sent over single TCP connection between client, server

Non-persistent HTTP Example

suppose user enters URL: www.comp.nus.edu.sg/~cs2105/demo.html (contains text, reference to a jpeg image)



Non-persistent HTTP Example



- ❖ This is an example of *non-persistent connection* (http 1.0).
 - One object per connection
- ❖ HTTP 1.1 allows *persistent connection* (to discuss).
 - Multiple objects per connection


Non-persistent HTTP: Response Time

RTT: time for a packet to travel from client to server and go back

HTTP response time:

- ❖ one RTT to establish TCP connection
- ❖ one RTT for HTTP request and the first few bytes of HTTP response to return
- ❖ file transmission time
- ❖ non-persistent HTTP response time =

$$2 * \text{RTT} + \text{file transmission time}$$



Persistent HTTP

non-persistent HTTP

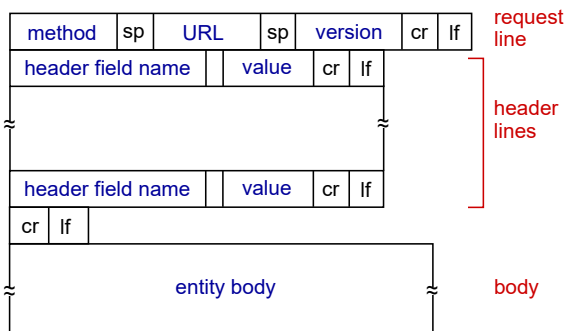
issues:

- ❖ requires 2 RTTs per object
- ❖ OS overhead for *each* TCP connection
- ❖ browsers often open parallel TCP connections to fetch referenced objects

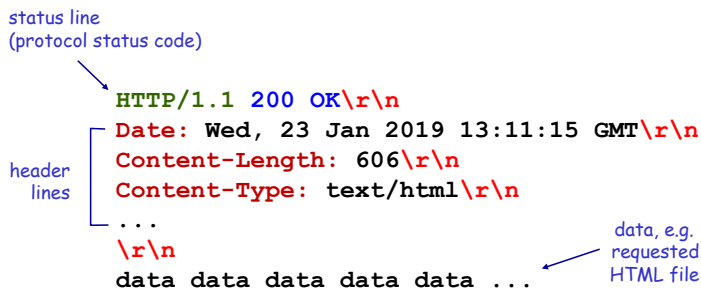
persistent HTTP:

- ❖ server leaves connection open after sending response
- ❖ subsequent HTTP messages between same client/server sent over the same TCP connection
- ❖ moreover, client may send requests as soon as it encounters a referenced object (**persistent with pipelining**)
- ❖ as little as one RTT for all the referenced objects

HTTP Request Message: General Format



Example HTTP Response Message



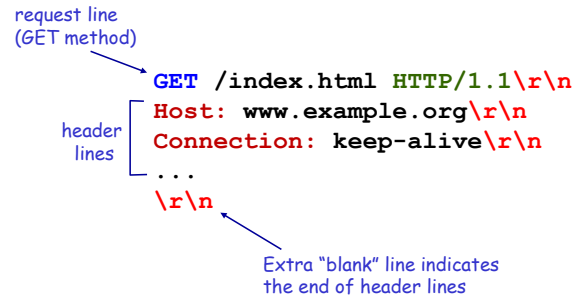
For a full list of request/response header fields, check <http://www.w3.org/Protocols/rfc2616/rfc2616-sec14.html>

Cookies

- ❖ HTTP is designed to be “stateless”.
 - Server maintains no information about past client requests.
- ❖ Sometimes it's good to maintain states (history) at server/client over multiple transactions.
 - E.g. shopping carts
- ❖ Cookie: http messages carry “state”
 - 1) cookie header field of HTTP *request* / *response* messages
 - 2) cookie file kept on user's host, managed by user's browser
 - 3) back-end database at Web site

Example HTTP Request Message

- ❖ Two types of HTTP messages: *request*, *response*
- ❖ HTTP request message:



HTTP Request Method Types

HTTP/1.0:

- ❖ GET
 - web page often includes form input
 - input is uploaded to server in entity body
- ❖ POST
 - asks server to leave requested object out of response

HTTP/1.1:

- ❖ GET, POST, HEAD
- ❖ PUT
 - uploads file in entity body to path specified in URL field
- ❖ DELETE
 - deletes file specified in the URL field

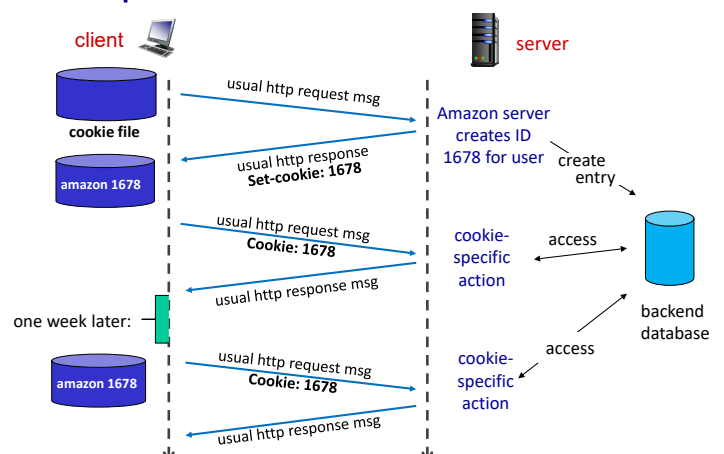
HTTP Response Status Codes

- ❖ Status code appears in 1st line in server-to-client response message.
- ❖ Some sample codes:

- 200 OK**
 - request succeeded, requested object later in this msg
- 301 Moved Permanently**
 - requested object moved, new location specified later in this msg (Location:)
- 403 Forbidden**
 - server declines to show the requested webpage
- 404 Not Found**
 - requested document not found on this server


For a full list of status codes, check <http://www.w3.org/Protocols/rfc2616/rfc2616-sec10.html>

Keep User States with Cookie



Conditional GET

- ❖ **Goal:** don't send object if (client) cache has up-to-date cached version
- ❖ **cache:** specify date of cached copy in HTTP request
`If-modified-since: <date>`
- ❖ **server:** response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not Modified`



Lectures 2&3: Roadmap

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- 2.4 DNS
- 2.7 Socket programming

Domain Name System [RFC 1034, 1035]

- ❖ Two ways to identify a host:
 - **Hostname**, e.g., `www.example.org`
 - **IP address**, e.g., `93.184.216.34`
- ❖ **DNS (Domain Name System)** translates between the two.
 - A client must carry out a DNS query to determine the IP address corresponding to the server name (e.g., `www.example.org`) prior to the connection.

DNS: Resource Records (RR)

- ❖ Mapping between host names and IP addresses (and others) are stored as **resource records (RR)**.

RR format: (name, value, type, ttl)

type = A

- **name** is hostname
- **value** is IP address

type = CNAME

- **name** is alias name for some "canonical" (the real) name
- **value** is canonical name
- e.g. `www.nus.edu.sg` is really `mgnzsqc.x.incapdns.net`

type = NS


- **name** is domain (e.g., `nus.edu.sg`)
- **value** is hostname of authoritative name server for this domain

type = MX

- **value** is name of mail server associated with **name**

Distributed, Hierarchical Database

- ❖ DNS stored RR in distributed databases implemented in hierarchy of many name servers.




A client wants IP address for `www.facebook.com`:

- ❖ client queries root server to find `.com` DNS server
- ❖ client queries `.com` DNS server to get `facebook.com` DNS server
- ❖ client queries `facebook.com` DNS server to get IP address for `www.facebook.com`

Root Servers

- ❖ Answers requests for records in the root zone by returning a list of the authoritative name servers for the appropriate top-level domain (TLD).



13 root name "servers" worldwide

TLD and Authoritative Servers

Top-level domain (TLD) servers:

- ❖ responsible for `com`, `org`, `net`, `edu`, ... and all top-level country domains, e.g., `uk`, `sg`, `jp`

Authoritative servers:

- ❖ Organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts (e.g. Web, mail)
- ❖ can be maintained by organization or service provider


Local DNS Server

- ❖ Does not strictly belong to hierarchy
- ❖ Each ISP (residential ISP, company, university) has one local DNS server.
 - also called "default name server"
- ❖ When host makes a DNS query, query is sent to its local DNS server
 - Retrieve name-to-address translation from local cache
 - Local DNS server acts as proxy and forwards query into hierarchy if answer is not found locally

DNS Caching

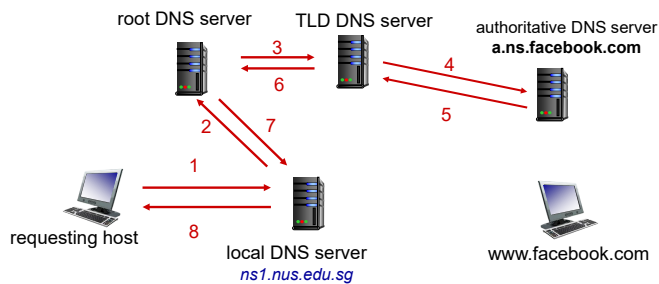
- ❖ Once a name server learns mapping, it *caches* mapping.
 - cached entries may be out-of-date (best effort name-to-address translation!)
 - cached entries expire after some time (TTL).
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- ❖ Update/notify mechanisms proposed IETF standard
 - RFC 2136
- ❖ DNS runs over **UDP**.

DNS Name Resolution



- ❖ This is known as *iterative query*.

DNS Name Resolution



- ❖ This is known as *recursive query*.
 - rarely happens in practice

Lectures 2&3: Roadmap

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- 2.4 DNS
- 2.7 Socket programming

Processes


- ❖ Process: program running within a host.
 - Within the same host, two processes communicate using *inter-process communication* (defined by OS).
 - Processes in different hosts communicate by exchanging *messages* (according to protocols).

In C/S model
server process waits to be contacted
client process initiates the communication

Addressing Processes

- ❖ **IP address** is used to identify a host device
 - A 32-bit integer (e.g. 137.132.21.27)
- ❖ **Question:** is IP address of a host suffice to identify a process running inside that host?

A: no, many processes may run concurrently in a host.



Addressing Processes

- ❖ A process is identified by (**IP address**, **port number**).
 - Port number is 16-bit integer (1-1023 are reserved for standard use).
- ❖ Example port numbers
 - HTTP server: 80
 - SMTP server: 25
- ❖ IANA coordinates the assignment of port number:
 - <http://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml>

Analogy

Postal service:


- ❖ *deliver letter to the doorstep*: home address
- ❖ *dispatch letter to the right person in the house*: name of the receiver as stated on the letter

Protocol service:

- ❖ *deliver packet to the right host*: IP address of the host
- ❖ *dispatch packet to the right process in the host*: port number of the process

Sockets


- ❖ **Socket** is the software interface between app processes and transport layer protocols.
 - Process sends/receives messages to/from its **socket**.
 - Programming-wise: a set of **APIs**



Socket Programming with UDP

UDP: no “connection” between client and server

- ❖ Sender (client) explicitly attaches destination IP address and port number to **each packet**.
- ❖ Receiver (server) extracts sender IP address and port number from the received packet.



Example: UDP Echo Server

```
from socket import *  # ← include Python's socket library

serverPort = 2105

# create a socket
serverSocket = socket(AF_INET, SOCK_DGRAM)

# bind socket to local port number 2105
serverSocket.bind(('', serverPort))
print('Server is ready to receive message')

# extract client address from received packet
message, clientAddress = serverSocket.recvfrom(2048)

serverSocket.sendto(message, clientAddress)


serverSocket.close()
```

IPv4 UDP socket

receive datagram buffer size: 2048B

Socket Programming with TCP

- ❖ When client creates socket, client TCP establishes a connection to server TCP.
- ❖ When contacted by client, **server TCP creates a new socket** for server process to communicate with that client.
 - allows server to talk with multiple clients individually.



Socket Programming



- ❖ Applications (or processes) treat the Internet as a black box, sending and receiving messages through sockets.
- ❖ Two types of sockets
 - **TCP**: reliable, byte stream-oriented socket
 - **UDP**: unreliable datagram socket
- ❖ Now let's write a simple client/server application that **client sends a line of text to server, and server echoes it**.
 - We will demo both **TCP socket version** and **UDP socket version**.

UDP: Client/server Socket Interaction

Server (running on serverIP)

```
create serverSocket,
port = x:

read datagram from
serverSocket

write reply to serverSocket
specifying client address,
port number
```

Client

```
create clientSocket

create datagram with serverIP
and port = x; send datagram via
clientSocket

read datagram from
clientSocket

close clientSocket
```

Example: UDP Echo Client

```
from socket import *

serverName = 'localhost' # client, server on the same host
serverPort = 2105

clientSocket = socket(AF_INET, SOCK_DGRAM)

message = input('Enter a message: ')

# send msg to server address
clientSocket.sendto(message.encode(), (serverName, serverPort))

receivedMsg, serverAddress = clientSocket.recvfrom(2048)

print('from server:', receivedMsg.decode())

clientSocket.close()
```

convert message from string to byte and send it

convert from byte to string

TCP: Client/server Socket Interaction

Server (running on serverIP)

```
create serverSocket, port = x

wait for incoming
connection request
connectionSocket

read request from
connectionSocket

write reply to
connectionSocket

close connectionSocket
```

Client

```
create clientSocket,
connect to serverIP, port = x

send request using clientSocket

read reply from clientSocket

close clientSocket
```

Example: TCP Echo Server

```
from socket import *

serverPort = 2105

serverSocket = socket(AF_INET, SOCK_STREAM)

serverSocket.bind(('', serverPort))

serverSocket.listen(1)
print('Server is ready to receive message')

connectionSocket, clientAddr = serverSocket.accept()
message = connectionSocket.recv(2048)

connectionSocket.send(message)

connectionSocket.close()
```

Annotations for TCP Echo Server:

- `serverSocket = socket(AF_INET, SOCK_STREAM)`: TCP socket
- `serverSocket.listen(1)`: listens for incoming TCP request (not available in UDP socket)
- `connectionSocket, clientAddr = serverSocket.accept()`: returns a **new** socket to communicate with client socket

Example: TCP Echo Client

```
from socket import *

serverName = 'localhost'
serverPort = 2105

clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))

message = input('Enter a message: ')

clientSocket.send(message.encode())

receivedMsg = clientSocket.recv(2048)

print('from server:', receivedMsg.decode())

clientSocket.close()
```

Annotations for TCP Echo Client:

- `clientSocket.connect((serverName, serverPort))`: establish a connection
- `clientSocket.send(message.encode())`: no need to attach server name, port

TCP Socket vs. UDP Socket

- ❖ In TCP, two processes communicate as if there is a pipe between them. The pipe remains in place until one of the two processes closes it.
 - When one of the processes wants to send more bytes to the other process, it simply writes data to that pipe.
 - The sending process doesn't need to attach a destination IP address and port number to the bytes in each sending attempt as the logical pipe has been established (which is also reliable).
- ❖ In UDP, programmers need to form UDP datagram packets explicitly and attach destination IP address / port number to every packet.

Lectures 2&3: Summary

- ❖ Application architectures
 - Client-server
 - P2P
- ❖ Application service requirements:
 - reliability, throughput, delay, security
- ❖ Specific protocols:
 - HTTP
 - DNS
- ❖ Internet transport service model
 - connection-oriented, reliable: TCP
 - Connection-less, unreliable: UDP

Lectures 2&3: Summary

- ❖ Socket programming
 - **TCP socket**
 - When contacted by client, server TCP creates new socket.
 - Server uses (client IP + port #) to distinguish clients.
 - When client creates its socket, client TCP establishes connection to server TCP.
 - **UDP socket**
 - Server uses **one socket** to serve all clients.
 - No connection is established before sending data.
 - Sender explicitly attaches **destination IP address** and **port #** to **each packet**.
 - Transmitted data may be lost or received out-of-order.